



LFS258

Kubernetes Fundamentals

Version 2018-01-16



© Copyright the Linux Foundation 2018. All rights reserved.

The training materials provided or developed by The Linux Foundation in connection with the training services are protected by copyright and other intellectual property rights.

Open source code incorporated herein may have other copyright holders and is used pursuant to the applicable open source license.

The training materials are provided for individual use by participants in the form in which they are provided. They may not be copied, modified, distributed to non-participants or used to provide training to others without the prior written consent of The Linux Foundation.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the **Linux Foundation**

<http://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please contact:

training@linuxfoundation.org

Contents

- 1 Introduction 1**
 - 1.1 Labs 1
- 2 Basics of Kubernetes 3**
 - 2.1 Labs 3
- 3 Kubernetes Architecture 5**
 - 3.1 Labs 5
- 4 Installation and Configuration 7**
 - 4.1 Labs 7
- 5 APIs and Access 27**
 - 5.1 Labs 27
- 6 Managing State With Deployments 33**
 - 6.1 Labs 33
- 7 Volumes and Data 41**
 - 7.1 Labs 41
- 8 Services 55**
 - 8.1 Labs 55
- 9 Ingress 61**
 - 9.1 Labs 61
- 10 API Objects 63**
 - 10.1 Labs 63
- 11 Scheduling 67**
 - 11.1 Labs 67
- 12 Logging and Troubleshooting 75**
 - 12.1 Labs 75
- 13 Custom Resource Definition 79**
 - 13.1 Labs 79
- 14 Kubernetes Federation 83**

14.1 Labs	83
15 Helm	85
15.1 Labs	85
16 Security	89
16.1 Labs	89

List of Figures

4.1 External Access via Browser 18

Chapter 1

Introduction



1.1 Labs

Exercise 1.1: Configuring the System for **sudo**

It is very dangerous to run a **root shell** unless absolutely necessary: a single typo or other mistake can cause serious (even fatal) damage.

Thus, the sensible procedure is to configure things such that single commands may be run with superuser privilege, by using the **sudo** mechanism. With **sudo** the user only needs to know their own password and never needs to know the root password.

If you are using a distribution such as **Ubuntu**, you may not need to do this lab to get **sudo** configured properly for the course. However, you should still make sure you understand the procedure.

To check if your system is already configured to let the user account you are using run **sudo**, just do a simple command like:

```
$ sudo ls
```

You should be prompted for your user password and then the command should execute. If instead, you get an error message you need to execute the following procedure.

Launch a root shell by typing **su** and then giving the **root** password, not your user password.

On all recent **Linux** distributions you should navigate to the `/etc/sudoers.d` subdirectory and create a file, usually with the name of the user to whom root wishes to grant **sudo** access. However, this convention is not actually necessary as **sudo** will scan all files in this directory as needed. The file can simply contain:

```
student ALL=(ALL) ALL
```

if the user is student.

An older practice (which certainly still works) is to add such a line at the end of the file `/etc/sudoers`. It is best to do so using the **visudo** program, which is careful about making sure you use the right syntax in your edit.

You probably also need to set proper permissions on the file by typing:

```
$ chmod 440 /etc/sudoers.d/student
```

(Note some **Linux** distributions may require 400 instead of 440 for the permissions.)

After you have done these steps, exit the root shell by typing `exit` and then try to do `sudo ls` again.

There are many other ways an administrator can configure **sudo**, including specifying only certain permissions for certain users, limiting searched paths etc. The `/etc/sudoers` file is very well self-documented.

However, there is one more setting we highly recommend you do, even if your system already has **sudo** configured. Most distributions establish a different path for finding executables for normal users as compared to root users. In particular the directories `/sbin` and `/usr/sbin` are not searched, since **sudo** inherits the `PATH` of the user, not the full root user.

Thus, in this course we would have to be constantly reminding you of the full path to many system administration utilities; any enhancement to security is probably not worth the extra typing and figuring out which directories these programs are in. Consequently, we suggest you add the following line to the `.bashrc` file in your home directory:

```
PATH=$PATH:/usr/sbin:/sbin
```

If you log out and then log in again (you don't have to reboot) this will be fully effective.

Chapter 2

Basics of Kubernetes



2.1 Labs

Exercise 2.1: View Online Resources

Visit kubernetes.io

With such a fast changing project, it is important to keep track of updates. The main place to find documentation of the current version is <https://kubernetes.io/>.

1. Open a browser and visit the <https://kubernetes.io/> website.
2. In the upper right hand corner, use the drop down to view the versions available. It will say something like v1.8.
3. Select the top level link for Documentation. The links on the left of the page can be helpful in navigation.
4. As time permits navigate around other sub-pages such as SETUP, CONCEPTS, and TASKS to become familiar to the layout.

Track Kubernetes Issues

There are hundreds, perhaps thousands, working on Kubernetes every day. With that many people working in parallel there are good resources to see if others are experiencing a similar outage. Both the source code as well as feature and issue tracking are currently on github.com.

1. To view the main page use your browser to visit <https://github.com/kubernetes/kubernetes/>
2. Click on various sub-directories and view the basic information available.
3. Update your URL to point to <https://github.com/kubernetes/kubernetes/issues>. You should see a series of issues, feature requests, and support communication.

4. In the search box you probably see some existing text like `is:issue is:open` which allows you to filter on the kind of information you would like to see. Append the search string to read: `is:issue is:open label:kind/bug` then press enter.
5. You should now see bugs in descending date order. Across the top of the issues a menu area allows you to view entries by author, labels, projects, milestones, and assignee as well. Take a moment to view the various other selection criteria.
6. Some times you may want to exclude a kind of output. Update the URL again, but precede the label with a minus sign, like: `is:issue is:open -label:kind/bug`. Now you see everything except bug reports.
7. Explore the page with the remaining time left.

Chapter 3

Kubernetes Architecture



3.1 Labs

- The lab for this chapter will appear in the next section, after we have discussed installation.

Chapter 4

Installation and Configuration



4.1 Labs

Exercise 4.1: Install Kubernetes

Overview

There are several Kubernetes installation tools provided by various vendors. In this lab we will learn to use **kubeadm** which is a **Beta**-release product. As an independent tool, it is planned to become the primary manner to build a Kubernetes cluster.

The labs were written using **Ubuntu** instances running on **Google Cloud Platform (GCP)**. They have been written to be vendor-agnostic so could run on AWS, local hardware, or inside of virtualization to give you the most flexibility and options. Each platform will have different access methods and considerations.

If using your own equipment you should configure **sudo** access as shown in a previous lab. While most commands are run as a regular user, there are some which require root privilege. If you are accessing the nodes remotely, such as with **GCP** or **AWS**, you will need to use an SSH client such as a local terminal or **PuTTY** if not using **Linux** or a Mac. You can download **PuTTY** from www.putty.org. You would also require a **.pem** or **.ppk** file to access the nodes. Each cloud provider will have a process to download or create this file. If attending in-person instructor led training the file will be made available during class.

In the following exercise we will install Kubernetes on a single node then grow our cluster, adding more compute resources. Both nodes used are the same size, providing 2 vCPUs and 7.5G of memory. Smaller nodes could be used, but would run slower.

Various exercises will use YAML files, which are included in the text. You are encouraged to write the files when possible, as the syntax of YAML has whitespace indentation requirements that are important to learn. The files have also been made available as a compressed **tar** file. You can view the resources by navigating to this URL:

<https://training.linuxfoundation.org/cm/LFS258/>

To login use user: LFtraining and a password of: Penguin2014

Once you find the name and link of the current file, which will change as the course updates, use **wget** to download the file into your node from the command line then expand it like this:

```
$ wget \
https://training.linuxfoundation.org/cm/LFS258/LFS258_V2018-01-08_SOLUTIONS.tar.bz2 \
--user=LFtraining --password=Penguin2014

$ tar -xvf LFS258_V2018-01-08_SOLUTIONS.tar.bz2
```

Install Kubernetes

Log into your nodes. If attending in-person instructor led training the node IP addresses will be provided by the instructor. You will need to use a `.pem` or `.ppk` key for access, depending on if you are using **ssh** from a terminal or **PuTTY**. The instructor will provide this to you.

1. Open a terminal session on your first node. For example, connect via **PuTTY** or **SSH** session to the first **GCP** node. The user name may be different than the one shown, `student`. The IP used in the example will be different than the one you will use.

```
[student@laptop ~]$ ssh -i LFS458.pem student@35.226.100.87
The authenticity of host '54.214.214.156 (35.226.100.87)' can't be established.
ECDSA key fingerprint is SHA256:IPvznbkx93/Wc+ACwXrCcDDgvBwmvEXC9vmYhk2Wo1E.
ECDSA key fingerprint is MD5:d8:c9:4b:b0:b0:82:d3:95:08:08:4a:74:1b:f6:e1:9f.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '35.226.100.87' (ECDSA) to the list of known hosts.
<output_omitted>
```

2. Become root and update and upgrade the system. Answer any questions to use the defaults.

```
student@lfs458-node-1a0a:~$ sudo -i

root@lfs458-node-1a0a:~# apt-get update && apt-get upgrade -y
<output_omitted>
```

3. The main choices for a container environment are **Docker** and **CoreOS Rocket - rkt**. We will use **Docker** for class, as **rkt** requires a fair amount of extra work to enable for Kubernetes.

```
root@lfs458-node-1a0a:~# apt-get install -y docker.io
<output-omitted>
```

4. Add new repo for kubernetes. You could also get a tar file or use code from GitHub. Create the file and add an entry for the main repo for your distribution.

```
root@lfs458-node-1a0a:~# vim /etc/apt/sources.list.d/kubernetes.list
deb http://apt.kubernetes.io/ kubernetes-xenial main
```

5. Add a GPG key for the packages.

```
root@lfs458-node-1a0a:~# curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg \
| apt-key add -
OK
```

6. Update with new repo.

```
root@lfs458-node-1a0a:~# apt-get update
<output-omitted>
```

Install the software. There are regular releases, the newest of which can be used by omitting the equal and version information on the command line. Historically new version have lots of changes and a good chance of a bug or two.

7. `root@lfs458-node-1a0a:~# apt-get install -y kubeadm=1.9.1-00 kubectl=1.9.1-00`
 <output-omitted>

8. Deciding which pod network to use for Container Networking Interface (**CNI**), should take into account the expected demands on the cluster. There can be only one pod network per cluster, although the **CNI-Genie** project is trying to change this.

The network must allow container to container, pod-to-pod, pod-to-service, and external-to-service communications. As **Docker** uses host-private networking, using the `docker0` virtual bridge and `veth` interfaces you would need to be on that host to communicate. **Flannel** is maintained by **CoreOS**. **Project Calico**, **OVN**, **Contrails**, and **OVS** are some of the several other options.

Download the **Flannel** file. Should you want to use **Calico**, those directions follow.

```
root@lfs458-node-1a0a:~# wget \
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

9. Take a look at some of the settings and network configurations. After taking a look at the file, determine the network address **Flannel** will expect. We will use this when we initialize the master server.

```
root@lfs458-node-1a0a:~# less kube-flannel.yml

root@lfs458-node-1a0a:~# grep Network kube-flannel.yml
  "Network": "10.244.0.0/16",
  hostNetwork: true
```

10. **ONLY IF YOU ARE USING CALICO, NOT FLANNEL** download the configuration file for **Calico**. Once downloaded look for the expected IP range for containers. It is different than **Flannel**. A short url is shown, for this URL: <https://docs.projectcalico.org/v2.6/getting-started/kubernetes/installation/hosted/kubeadm/1.6/calico.yaml>

```
root@lfs458-node-1a0a:~# wget https://goo.gl/eWLkzb -O calico.yaml

root@lfs458-node-1a0a:~# less calico.yaml
....
    # Configure the IP Pool from which Pod IPs will be chosen.
    - name: CALICO_IPV4POOL_CIDR
      value: "192.168.0.0/16"
....
```

11. Initialize the master. Read through the output line by line. Note that the software is in beta as well as some of the differences expected in future versions. At the end are directions to run as a non-root user. The token is mentioned as well. This information can be found later with the **kubeadm token list** command. The output also directs you to create a pod network to the cluster, which will be our next step. Pass the network settings **Flannel** will expect to find.

```
root@lfs458-node-1a0a:~# kubeadm init --pod-network-cidr 10.244.0.0/16
[kubeadm] WARNING: kubeadm is in beta, please do not use it for production clusters.
[init] Using Kubernetes version: v1.9.1
[init] Using Authorization modes: [Node RBAC]
[preflight] Running pre-flight checks
```

<output-omitted>

Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run (as a regular user):

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<http://kubernetes.io/docs/admin/addons/>

You can now join any number of machines by running the following on each node as root:

```
kubeadm join --token 563c3c.9c978c8c0e5fbb4 10.128.0.3:6443
--discovery-token-ca-cert-hash sha256:726e98586a8d12d428c0ee46
cbea90c094b8a78cb272917e2681f7b75abf875f
```

12. Follow the suggestion to allow a non-root user access to the cluster. Take a quick look at the configuration file once it has been copied and the permissions fixed.

```
root@lfs458-node-1a0a:~# exit
logout
student@lfs458-node-1a0a:~$ mkdir -p $HOME/.kube

student@lfs458-node-1a0a:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config

student@lfs458-node-1a0a:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config

student@lfs458-node-1a0a:~$ less .kube/config
apiVersion: v1
clusters:
- cluster:
<output_omitted>
```

13. Apply the configuration to your cluster. Remember to copy the file to the current, non-root user directory first. When it finished you should see a new flannel interface. It may take up to a minute to be created.

```
student@lfs458-node-1a0a:~$ sudo cp /root/kube-flannel.yml .

student@lfs458-node-1a0a:~$ kubectl apply -f kube-flannel.yml
clusterrole "flannel" created
clusterrolebinding "flannel" created
serviceaccount "flannel" created
configmap "kube-flannel-cfg" created
daemonset "kube-flannel-ds" created

student@lfs458-node-1a0a:~$ ip a
<output_omitted>
4: flannel.1: <BROADCAST,MULTICAST> mtu 8951 qdisc noop state DOWN group default
    link/ether 32:44:47:b7:78:85 brd ff:ff:ff:ff:ff:ff
```

14. View the available nodes of the cluster. It can take a minute or two for the status to change from NotReady to Ready. The NAME field can be used to look at the details. Your node name will be different.

```
student@lfs458-node-1a0a:~$ kubectl get node
NAME                STATUS    AGE           VERSION
lfs458-node-1a0a    Ready    1m            v1.9.1
```

15. Look at the details of the node. Work line by line to view the resources and their current status. Notice the status of Taints. The master won't allow pods by default for security reasons. Take a moment to read each line of output, some appear to be an error until you notice the status shows False.

```
student@lfs458-node-1a0a:~$ kubectl describe node lfs458-node-1a0a
Name:                lfs458-node-1a0a
Role:
Labels:              beta.kubernetes.io/arch=amd64
                    beta.kubernetes.io/os=linux
                    kubernetes.io/hostname=lfs458-node-1a0a
                    node-role.kubernetes.io/master=
Annotations:         node.alpha.kubernetes.io/ttl=0
                    volumes.kubernetes.io/controller-managed-attach-detach=true
Taints:              node-role.kubernetes.io/master:NoSchedule
<output_omitted>
```


16. Determine if the DNS and flannel pods are ready for use. They should all show a status of Running. It may take a minute or two to transition from Pending.

```
student@lfs458-node-1a0a:~$ kubectl get pods --all-namespaces
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  etcd-lfs458-node-1a0a                 1/1     Running   0          12m
kube-system  kube-apiserver-lfs458-node-1a0a       1/1     Running   0          12m
kube-system  kube-controller-manager-lfs458-node-1a0a 1/1     Running   0          12m
kube-system  kube-dns-2425271678-w80vx            3/3     Running   0          13m
kube-system  kube-flannel-ds-wj92l                 1/1     Running   0          1m
kube-system  kube-proxy-5st9z                     1/1     Running   0          13m
kube-system  kube-scheduler-lfs458-node-1a0a       1/1     Running   0          12m
```

17. Allow the master server to run other pods. Note the minus sign at the end, which is the syntax to remove a taint.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes --all node-role.kubernetes.io/master-
node "lfs458-node-1a0a" untainted

student@lfs458-node-1a0a:~$ kubectl describe node lfs458-node-1a0a | grep -i taint
Taints:                                     <none>
```

18. While many objects have short names, a **kubectl** command can be a lot to type. We will enable **bash** auto-completion. Begin by adding the settings to the current shell. Then update the `~/.bashrc` file to make it persistent.

```
student@lfs458-node-1a0a:~$ source <(kubectl completion bash)

student@lfs458-node-1a0a:~$ echo "source <(kubectl completion bash)" >> ~/.bashrc
```

19. Test by describing the node again. Type the first three letters of the sub-command then type the Tab key.

```
student@lfs458-node-1a0a:~$ kubectl des<Tab> n<Tab><Tab> lfs458-<Tab>
```

Exercise 4.2: Grow the Cluster

Open another terminal and connect into a your second node. Install **Docker** and Kubernetes software. These are the same steps we did on the master node.

1. Using the same process as before connect to a second node. If attending ILT use the same `.pem` key and a new IP provided by the instructor to access the new node. Giving a title or color to the new terminal window is probably a good idea to keep track of the two systems. The prompts can look very similar.

```
student@lfs458-node-2b2b:~$ sudo -i

root@lfs458-node-2b2b:~# apt-get update && apt-get upgrade -y

root@lfs458-node-2b2b:~# apt-get install -y docker.io

root@lfs458-node-2b2b:~# vim /etc/apt/sources.list.d/kubernetes.list
deb http://apt.kubernetes.io/ kubernetes-xenial main

root@lfs458-node-2b2b:~# curl -s \
    https://packages.cloud.google.com/apt/doc/apt-key.gpg \
    | apt-key add -

root@lfs458-node-2b2b:~# apt-get update

root@lfs458-node-2b2b:~# apt-get install -y kubeadm=1.9.1-00 kubelet=1.9.1-00
```

2. Find the IP address of your master server. The interface name will be different depending on where the node is running. Currently inside of **GCE** the primary interface for this node type is `ens4`. Your interfaces names may be different. From the output we know our master node IP is `10.128.0.3`.

```
student@lfs458-node-1a0a:~$ ip addr show ens4 | grep inet
    inet 10.128.0.3/32 brd 10.128.0.3 scope global ens4
    inet6 fe80::4001:aff:fe8e:2/64 scope link
```

- Find the token on the master node. The token lasts 24 hours by default. If the it has been longer, and no token is present you can generate a new one with the **sudo kubeadm token create** command.

```
student@lfs458-node-1a0a:~$ sudo kubeadm token list
TOKEN          TTL      EXPIRES  USAGES  DESCRIPTION
27eee4.6e66ff60318da929  23h      2017-11-03T13:27:33Z
authentication,signing The default bootstrap token generated
by 'kubeadm init'....
```

- Starting in v1.9 you should create and use a Discovery Token CA Cert Hash created from the master to ensure the node joins the cluster in a secure manner. Run this on the master node or wherever you have a copy of the CA file. You will get a long string as output.

```
student@lfs458-node-1a0a:~$ openssl x509 -pubkey \
    -in /etc/kubernetes/pki/ca.crt | openssl rsa \
    -pubin -outform der 2>/dev/null | openssl dgst \
    -sha256 -hex | sed 's/^.* //'
6d541678b05652e1fa5d43908e75e67376e994c3483d6683f2a18673e5d2a1b0
```

- Use the token and hash, in this case as sha256:<hash> to join the cluster from the second node. Use the **private** IP address of the master server and port 6443. The output of the **kubeadm init** on the master also has an example to use, should it still be available.

```
root@lfs458-node-2b2b:~# kubeadm join \
    --token 27eee4.6e66ff60318da929 10.128.0.3:6443
    --discovery-token-ca-cert-hash \
    sha256:6d541678b05652e1fa5d43908e75e67376e994c3483d6683f2a18673e5d2a1b0
```

```
[preflight] Running pre-flight checks.
[WARNING FileExisting-crictl]: crictl not found in system path
[discovery] Trying to connect to API Server "10.142.0.2:6443"
[discovery] Created cluster-info discovery client, requesting info from
"https://10.142.0.2:6443"
[discovery] Requesting info from "https://10.142.0.2:6443" again to
validate TLS against the pinned public key
[discovery] Cluster info signature and contents are valid and TLS
certificate validates against pinned roots, will
use API Server "10.142.0.2:6443"
[discovery] Successfully established connection with API Server
"10.142.0.2:6443"
This node has joined the cluster:
* Certificate signing request was sent to master and a response
  was received.
* The Kubelet was informed of the new secure connection details.
Run 'kubectl get nodes' on the master to see this node join the cluster.
```

- Try to run the **kubectl** command on the secondary system. It should fail. You do not have the cluster or authentication keys in your local **.kube/config** file.

```
root@lfs458-node-2b2b:~# exit

student@lfs458-node-2b2b:~$ kubectl get nodes
The connection to the server localhost:8080 was refused
- did you specify the right host or port?

student@lfs458-node-2b2b:~$ ls -l .kube
ls: cannot access '.kube': No such file or directory
```

- Verify the node has joined the cluster from the master node. You may need to wait a minute for the node to show a Ready state.

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
lfs458-node-2b2b    NotReady <none>    14s      v1.9.1
lfs458-node-1a0a    Ready     master    16m      v1.9.1
```

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
lfs458-node-2b2b    Ready     <none>    1m       v1.9.1
lfs458-node-1a0a    Ready     master    17m      v1.9.1
```

- View the current namespaces configured on the cluster.

```
student@lfs458-node-1a0a:~$ kubectl get namespace
NAME                STATUS    AGE
default             Active    17m
kube-public         Active    17m
kube-system         Active    17m
```

- View the networking on the master and second node. You should see a docker0, cni0, and a flannel.1 interface among others.

```
student@lfs458-node-1a0a:~$ ip a
<output_omitted>
```

- Start a new simple deployment from the command line. Verify it is running.

```
student@lfs458-node-1a0a:~$ kubectl run nginx --image nginx
deployment "nginx" created
```

```
student@lfs458-node-1a0a:~$ kubectl get deployments
NAME                DESIRED    CURRENT    UP-TO-DATE    AVAILABLE    AGE
nginx               1          1          1             1            6s
```

- View the details of the deployment. Remember auto-completion will work for sub-commands and resources as well.

```
student@lfs458-node-1a0a:~$ kubectl describe deployment nginx
Name:                nginx
Namespace:           default
CreationTimestamp:    Tue, 26 Sep 2017 21:49:51 +0000
Labels:               run=nginx
Annotations:          deployment.kubernetes.io/revision=1
Selector:             run=nginx
Replicas:             1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:         RollingUpdate
MinReadySeconds:      0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
<output_omitted>
```

- View the basic steps the cluster took in order to pull and deploy the new application. You should see about ten long lines of output.

```
student@lfs458-node-1a0a:~$ kubectl get events
<output_omitted>
```

- You can also view the output in **yaml** format, which could be used to create this deployment again or new deployments. Get the information but change the output to yaml. Note that halfway down there is status information of the current deployment.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx -o yaml
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  creationTimestamp: 2017-09-27T18:21:25Z
<output_omitted>
```

14. Run the command again and redirect the output to a file. Then edit the file. Remove the `creationTimestamp`, `resourceVersion`, `selfLink`, and `uid` lines. Also remove all the lines including and after `status:`, which should be somewhere around line 39, if others have already been removed.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx -o yaml > first.yaml

student@lfs458-node-1a0a:~$ vim first.yaml
```

15. Delete the existing deployment.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment nginx
deployment "nginx" deleted
```

16. Create the deployment again this time using the file.

```
student@lfs458-node-1a0a:~$ kubectl create -f first.yaml
deployment "nginx" created
```

17. Look at the yaml output of this iteration and compare it against the first. The time stamp, resource version and uid we had deleted are in the new file. These are generated for each resource we create, so we need to delete them from yaml files to avoid conflicts or false information. `status` should not be hard-coded either.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx -o yaml > second.yaml

student@lfs458-node-1a0a:~$ diff first.yaml second.yaml
<output_omitted>
```

18. The newly deployed **nginx** container is a light weight web server. We will need to create a `service` to view the default welcome page. Begin by looking at the help output. Note that there are several examples given, about halfway through the output.

```
student@lfs458-node-1a0a:~$ kubectl expose -h
<output_omitted>
```

19. Now try to gain access to the web server. As we have not declared a port to use you will receive an error.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment/nginx
error: couldn't find port via --port flag or introspection
See 'kubectl expose -h' for help and examples.
```

20. To change an existing configuration in a cluster can be done with subcommands `apply`, `edit` or `patch` for non-disruptive updates. The `apply` command does a three-way diff of previous, current, and supplied input to determine modifications to make. Changes not mentioned are unaffected. The `edit` function performs a `get`, opens an editor then an `apply` for you. You can update API objects in place with `JSON patch` and `merge patch` or `strategic merge patch` functionality.

If the configuration has resource fields which cannot be updated once initialized then a disruptive update would could be done using the `replace --force` option. This deletes first then re-creates a resource.

Edit the file. Find the container name, somewhere around line 31 and add the port information as shown below.

```
student@lfs458-node-1a0a:~$ vim first.yaml
.
spec:
  containers:
  - image: nginx
```

```

imagePullPolicy: Always
name: nginx
ports:
- containerPort: 80
  protocol: TCP
resources: {}

```

21. Apply the changes to the running deployment. You may get a warning that apply should be used on resources which were created with particular options. The command should still work.

```

student@lfs458-node-1a0a:~$ kubectl apply -f first.yaml
deployment "nginx" created

```

22. View the Pod and Deployment. Note the Pod was re-created.

```

student@lfs458-node-1a0a:~$ kubectl get deploy,pod
NAME          DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
deploy/nginx  1        1        1           1          8s

NAME          READY   STATUS    RESTARTS  AGE
nginx-7cbc4b4d9c-l8cgl  1/1     Running   0         8s

```

23. Try to expose the resource again.

```

student@lfs458-node-1a0a:~$ kubectl expose deployment/nginx
service "nginx" exposed

```

24. Verify the service configuration. First look at the service information, then at the endpoint information. Note the Cluster IP is not the current endpoint. Take note of the current endpoint IP. In the example below it is 10.244.1.99:80. We will use this information in a few steps.

```

student@lfs458-node-1a0a:~$ kubectl get svc nginx
NAME      CLUSTER-IP      EXTERNAL-IP  PORT(S)  AGE
nginx     10.100.61.122   <none>       80/TCP   3m

student@lfs458-node-1a0a:~$ kubectl get ep nginx
NAME      ENDPOINTS          AGE
nginx     10.244.1.99:80    4m

```

25. Determine which node the container is running on. Log into that node and use **tcpdump** to view traffic on the `flannel.1` interface. The second node in this example. Yours may be different. Leave that command running while you run **curl** in the following step. You should see several messages go back and forth, including a `HTTP: HTTP/1.1 200 OK` and a ack response to the same sequence.

```

student@lfs458-node-1a0a:~$ kubectl describe pod nginx-7cbc4b4d9c-d27xw \
| grep Node:
Node:    lfs458-node-2b2b/10.128.0.5

student@lfs458-node-2b2b:~$ sudo tcpdump -i flannel.1
tcpdump: verbose output suppressed, use -v or -vv for full protocol...
listening on flannel.1, link-type EN10MB (Ethernet), capture size...
<output_omitted>

```

26. Test access to the Cluster IP, port 80. You should see the generic `nginx` installed and working page. The output should be the same when you look at the `ENDPOINTS` IP address. If the **curl** command times out the pod may be running on the other node. Run the same command on that node and it should work.

```

student@lfs458-node-1a0a:~$ curl 10.100.61.122:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>

```

27. Stop the **tcpdump** command on the second terminal. Run **curl** again while using the **cni0** interface with **tcpdump** in a different terminal. You should see similar messages as on the **flannel.1** interfaces and even more about housekeeping. Again on the **veth** interface, yours may be different, and you'll also see MAC information. Leave **tcpdump** running for the next several steps, to see traffic as the deployment is scaled.

```
student@lfs458-node-2b2b:~$ sudo tcpdump -i cni0
<output_omitted>

student@lfs458-node-2b2b:~$ sudo tcpdump -i vethfa2c0158
<output_omitted>
```

28. Now scale up the deployment from one to three web servers.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx     1         1         1            1           12m

student@lfs458-node-1a0a:~$ kubectl scale deployment nginx --replicas=3
deployment "nginx" scaled

student@lfs458-node-1a0a:~$ kubectl get deployment nginx
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx     3         3         3            3           12m
```

29. View the current endpoints. There now should be three.

```
student@lfs458-node-1a0a:~$ kubectl get ep nginx
NAME      ENDPOINTS                                     AGE
nginx     10.244.0.66:80,10.244.1.100:80,10.244.1.99:80 10m
```

30. Find the oldest pod of the **nginx** deployment and delete it. The Tab key can be helpful for the long names. Use the AGE field to determine which was running the longest. You will notice activity in the other terminal where **tcpdump** is running, when you delete the pod.

```
student@lfs458-node-1a0a:~$ kubectl get po -o wide
NAME                                READY    STATUS    RESTARTS   AGE   IP
nginx-1423793266-7f1qw             1/1     Running   0          14m   10.244.0.66
nginx-1423793266-8w2nk             1/1     Running   0          1m    10.244.1.100
nginx-1423793266-fbt4b             1/1     Running   0          1m    10.244.1.101

student@lfs458-node-1a0a:~$ kubectl delete po nginx-1423793266-7f1qw
pod "nginx-1423793266-7f1qw" deleted
```

31. Wait a minute or two then view the pods again. One should be newer than the others. In the following example two minutes instead of four. If your **tcpdump** was using the **veth** interface of that container it will error out.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                                READY    STATUS    RESTARTS   AGE
nginx-1423793266-13p69             1/1     Running   0          2m
nginx-1423793266-8w2nk             1/1     Running   0          4m
nginx-1423793266-fbt4b             1/1     Running   0          4m
```

32. View the endpoints again. The original endpoint IP is no longer in use. You can delete any of the pods and the service will forward traffic to the existing backend pods.

```
student@lfs458-node-1a0a:~$ kubectl get ep nginx
NAME      ENDPOINTS                                     AGE
nginx     10.244.0.66:80,10.244.1.100:80,10.244.1.101:80 15m
```

33. Test access to the web server again, using the **ClusterIP** address, then any of the endpoint IP addresses. Even though the endpoints have changed you still have access to the web server. This access is only from within the cluster.

```
student@lfs458-node-1a0a:~$ curl 10.100.61.122:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
<output_omitted>
```

Exercise 4.3: Access from Outside the Cluster

You can access a Service from outside the cluster using a DNS add-on or **vi** environment variables. We will use environment variables to gain access to a Pod.

1. Begin by getting a list of the pods.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                                READY    STATUS    RESTARTS   AGE
nginx-1423793266-13p69             1/1      Running   0           8m
nginx-1423793266-8w2nk             1/1      Running   0           10m
nginx-1423793266-fbt4b             1/1      Running   0           10m
```

2. Choose one of the pods and use the `exec` command to run **printenv** inside the pod. The following example uses the first pod listed above.

```
student@lfs458-node-1a0a:~$ kubectl exec nginx-1423793266-13p69 \
-- printenv |grep KUBERNETES
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_SERVICE_HOST=10.96.0.1
KUBERNETES_SERVICE_PORT=443
NGINX_SERVICE_HOST=10.100.61.122
NGINX_SERVICE_PORT=80
<output_omitted>
```

3. Find and then delete the existing service for **nginx**.

```
student@lfs458-node-1a0a:~$ kubectl get svc
NAME            CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes      10.96.0.1     <none>         443/TCP    2d
nginx           10.100.61.122 <none>         80/TCP     40m
```

4. Delete the service.

```
student@lfs458-node-1a0a:~$ kubectl delete svc nginx
service "nginx" deleted
```

5. Create the service again, but this time pass the `LoadBalancer` type. Check to see the status and note the external ports mentioned. The output will show the `External-IP` as pending. **AWS** and **GCE** have a quirk where it always shows as pending.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment nginx --type=LoadBalancer
service "nginx" exposed
```

```
student@lfs458-node-1a0a:~$ kubectl get svc
NAME            CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes      10.96.0.1     <none>         443/TCP    2d
nginx           10.104.249.102 <pending>      80:32753/TCP 2s
```

6. Open a browser on your local system, not the GCE node, and use the public IP of your node and port 32753, shown in the output above. If running the labs on a remote system like **AWS** or **GCE** the `CLUSTER-IPs` are internal. Use the public IP you used with SSH to gain access.



Figure 4.1: External Access via Browser

7. Scale the deployment to zero replicas. Then test the web page again. It should fail.

```
student@lfs458-node-1a0a:~$ kubectl scale deployment nginx --replicas=0
deployment "nginx" scaled
```

```
student@lfs458-node-1a0a:~$ kubectl get po
No resources found.
```

8. Scale the deployment up to two replicas. The web page should work again.

```
student@lfs458-node-1a0a:~$ kubectl scale deployment nginx --replicas=2
deployment "nginx" scaled
```

```
student@lfs458-node-1a0a:~$ kubectl get po
. NAME                                READY   STATUS    RESTARTS   AGE
nginx-1423793266-7x181                1/1     Running   0           1m
nginx-1423793266-s6vcz                1/1     Running   0           1m
```

9. Delete the deployment to recover system resources.

```
student@lfs458-node-1a0a:~$ kubectl delete deployments nginx
deployment "nginx" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl delete ep nginx
endpoints "nginx" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl delete svc nginx
service "nginx" deleted
```

Overview

We will continue working with our cluster, which we built in the previous lab. We will work with `resource limits`, more with namespaces and then a complex deployment which you can explore to further understand the architecture and relationships.

Exercise 4.4: Working with CPU and Memory Constraints

Use **SSH** or **PuTTY** to connect to the nodes you installed in the previous exercise. We will deploy an application called **stress** inside a container, and then use `resource limits` to constrain the resources the application has access to use.

1. Use a container called `stress`, which we will name `hog`, to generate load. Verify you have a container running.

```
student@lfs458-node-1a0a:~$ kubectl run hog --image vish/stress
deployment "hog" created
```



```
student@lfs458-node-1a0a:~$ kubectl get deployments
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hog            1         1         1            1          12s
```

2. Use the describe argument to view details, then view the output in YAML format. Note there are no settings limiting resource usage. Instead, there are empty curly brackets.

```
student@lfs458-node-1a0a:~$ kubectl describe deployment hog
Name:          hog
Namespace:     default
CreationTimestamp: Wed, 27 Sep 2017 20:09:57 +0000
Labels:        run=hog
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl get deployment hog -o yaml
apiVersion: extensions/v1beta1
kind: Deployment
Metadata:

<output_omitted>

  template:
    metadata:
      creationTimestamp: null
      labels:
        run: hog
    spec:
      containers:
      - image: vish/stress
        imagePullPolicy: Always
        name: hog
        resources: {}
        terminationMessagePath: /dev/termination-log
<output_omitted>
```

3. We will use the YAML output to create our own configuration file.

```
student@lfs458-node-1a0a:~$ kubectl get deployment hog -o yaml > hog.yaml
```

4. We will need to remove the status output, creationTimestamp and other settings, as we don't want to set unique generated parameters. We will also add in memory limits found below.

```
student@lfs458-node-1a0a:~$ vim hog.yaml
.
  imagePullPolicy: Always
  name: hog
  resources:          # Edit to remove {}
    limits:          # Add these 4 lines
      memory: "4Gi"
    requests:
      memory: "2500Mi"
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
....
```

5. Replace the deployment using the newly edited file.

```
student@lfs458-node-1a0a:~$ kubectl replace -f hog.yaml
```

6. Verify the change has been made. The deployment should now show resource limits.

```
student@lfs458-node-1a0a:~$ kubectl get deployment hog -o yaml |less
....
resources:
  limits:
    memory: 4Gi
  requests:
    memory: 2500Mi
  terminationMessagePath: /dev/termination-log
....
```

7. View the `stdio` of the `hog` container. Note how much memory has been allocated.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                                READY    STATUS    RESTARTS   AGE
hog-64cbfcc7cf-lwq66               1/1      Running   0           2m

student@lfs458-node-1a0a:~$ kubectl logs hog-64cbfcc7cf-lwq66
I1102 16:16:42.638972      1 main.go:26] Allocating "0" memory, in
    "4Ki" chunks, with a 1ms sleep between allocations
I1102 16:16:42.639064      1 main.go:29] Allocated "0" memory
```

8. Open a second and third terminal to access both master and second nodes. Run **top** to view resource usage. You should not see unusual resource usage at this point. The **dockerd** and **top** processes should be using about the same amount of resources. The **stress** command should not be using enough resources to show up.
9. Edit the `hog` configuration file and add arguments for **stress** to consume CPU and memory.

```
student@lfs458-node-1a0a:~$ vim hog.yaml
resources:
  limits:
    cpu: "1"
    memory: "4Gi"
  requests:
    cpu: "0.5"
    memory: "500Mi"
args:
- -cpus
- "2"
- -mem-total
- "950Mi"
- -mem-alloc-size
- "100Mi"
- -mem-alloc-sleep
- "1s"
```

10. Delete and recreate the deployment. You should see CPU usage almost immediately and memory allocation happen in 100M chunks allocated to the **stress** program. Check both nodes as the container could be deployed to either. The next step will help if you have errors.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment hog
deployment "hog" deleted

student@lfs458-node-1a0a:~$ kubectl apply -f hog.yaml
deployment "hog" created
```

11. Should the resources not show as used, there may have been an issue inside of the container. Kubernetes shows it as running, but the actual workload has failed. Or the container may have failed; for example if you were missing a parameter the container may panic and show the following output.

```
student@lfs458-node-1a0a:~$ kubectl get pod
NAME                                READY    STATUS    RESTARTS   AGE
hog-1985182137-5bz2w               0/1      Error     1           5s
```

```
student@lfs458-node-1a0a:~$ kubectl logs hog-1985182137-5bz2w
panic: cannot parse '150mi': unable to parse quantity's suffix

goroutine 1 [running]:
panic(0x5ff9a0, 0xc820014cb0)
    /usr/local/go/src/runtime/panic.go:481 +0x3e6
k8s.io/kubernetes/pkg/api/resource.MustParse(0x7ffe460c0e69, 0x5, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0)
    /usr/local/google/home/vishnuk/go/src/k8s.io/kubernetes/pkg/api/resource/quantity.go:134 +0x287
main.main()
    /usr/local/google/home/vishnuk/go/src/github.com/vishh/stress/main.go:24 +0x43
```

12. Here is an example of an improper parameter. The container is running, but not allocating memory. It should show the usage requested from the YAML file.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                                READY    STATUS    RESTARTS   AGE
hog-1603763060-x3vnn               1/1      Running   0           8s

student@lfs458-node-1a0a:~$ kubectl logs hog-1603763060-x3vnn
I0927 21:09:23.514921      1 main.go:26] Allocating "0" memory, in "4Ki" chunks, with a 1ms sleep \
                        between allocations
I0927 21:09:23.514984      1 main.go:39] Spawning a thread to consume CPU
I0927 21:09:23.514991      1 main.go:39] Spawning a thread to consume CPU
I0927 21:09:23.514997      1 main.go:29] Allocated "0" memory
```

Exercise 4.5: Resource Limits for a Namespace

The previous steps set limits for that particular deployment. You can also set limits on an entire namespace. We will create a new namespace and configure the **hog** deployment to run within. When set **hog** should not be able to use the previous amount of resources.

1. Begin by creating a new namespace called `low-usage-limit` and verify it exists.

```
student@lfs458-node-1a0a:~$ kubectl create namespace low-usage-limit
namespace "low-usage-limit" created

student@lfs458-node-1a0a:~$ kubectl get namespace
NAME                STATUS    AGE
default             Active    1h
kube-public         Active    1h
kube-system         Active    1h
low-usage-limit     Active    42s
```

2. Create a YAML file which limits CPU and memory usage. The kind to use is `LimitRange`.

```
student@lfs458-node-1a0a:~$ vim low-resource-range.yaml
```

3. Create the `LimitRange` object and assign it to the newly created namespace `low-usage-limit`

```
student@lfs458-node-1a0a:~$ kubectl create -f low-resource-range.yaml \
--namespace=low-usage-limit
limitrange "low-resource-range" created
```

4. Verify it works. Remember that every command needs a namespace and context to work. Defaults are used if not provided.

```
student@lfs458-node-1a0a:~$ kubectl get LimitRange
No resources found.

student@lfs458-node-1a0a:~$ kubectl get LimitRange --all-namespaces
NAMESPACE          NAME                AGE
low-usage-limit    low-resource-range  48s
```

5. Create a new deployment in the namespace.

```
student@lfs458-node-1a0a:~$ kubectl run limited-hog \
    --image vish/stress -n low-usage-limit
deployment "limited-hog" created
```

6. List the current deployments. Note `hog` continues to run in the default namespace.

```
student@lfs458-node-1a0a:~$ kubectl get deployments --all-namespaces
```

NAMESPACE	NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
default	hog	1	1	1	1	25m
kube-system	kube-dns	1	1	1	1	2d
low-usage-limit	limited-hog	1	1	1	1	1m

7. View all pods within the namespace.

```
student@lfs458-node-1a0a:~$ kubectl get pods -n low-usage-limit
```

NAME	READY	STATUS	RESTARTS	AGE
limited-hog-2556092078-wnpnv	1/1	Running	0	3m

8. Look at the details of the pod. You will note it has the settings inherited from the entire namespace. The use of shell completion may not work as it doesn't exist in the default namespace.

```
student@lfs458-node-1a0a:~$ kubectl get pod limited-hog-2556092078-wnpnv -n low-usage-limit -o yaml
<output_omitted>
spec:
  containers:
  - image: vish/stress
    imagePullPolicy: Always
    name: limited-hog
    resources:
      limits:
        cpu: "1"
        memory: 500Mi
      requests:
        cpu: 500m
        memory: 100Mi
    terminationMessagePath: /dev/termination-log
<output_omitted>
```

9. Copy and edit the config file for the original `hog` file. Edit the namespace: line only such that a new deployment would be in the `low-usage-limit` namespace.

```
student@lfs458-node-1a0a:~$ cp hog.yaml hog2.yaml
```

```
student@lfs458-node-1a0a:~$ vim hog2.yaml
```

```
....
  labels:
    run: hog
    name: hog
    namespace: low-usage-limit    #<<--- This line
spec:
....
```

10. Open up extra terminal sessions so you can have **top** running in each. When the new deployment is created it will probably be scheduled on the node not yet under any stress.

Create the deployment.

```
student@lfs458-node-1a0a:~$ kubectl create -f hog2.yaml
deployment "hog" created
```

11. View the deployments. Note there are two with the same name, but in different namespaces.

```
student@lfs458-node-1a0a:~$ kubectl get deployments --all-namespaces
NAMESPACE      NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
default        hog           1         1         1            1          15m
kube-system    kube-dns      1         1         1            1          3h
low-usage-limit hog           1         1         1            1          27s
low-usage-limit limited-hog    1         1         1            1          4m
```

- Look at the **top** output. You should find that both deployments are using the same amount of resources, once the memory is fully allocated. Per-deployment settings override the global namespace settings. You should see something like the following lines one from each node

```
25128 root      20   0 958532 954672   3180 R 100.0 11.7   0:52.27 stress
24875 root      20   0 958532 954800   3180 R 100.3 11.7  41:04.97 stress
```

- Delete the hog deployments to recover system resources.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment hog -n low-usage-limit
deployment "hog" deleted

student@lfs458-node-1a0a:~$ kubectl delete deployment hog
deployment "hog" deleted
```

Exercise 4.6: More Complex Deployment

We will now deploy a more complex demo application to test the cluster. When completed it will be a sock shopping site. The short URL is shown below for: <https://raw.githubusercontent.com/microservices-demo/microservices-demo/master/deploy/kubernetes/complete-demo.yaml>

- Begin by downloading the pre-made YAML file from github.

```
student@lfs458-node-1a0a:~$ wget https://tinyurl.com/y8bn2awp -O complete-demo.yaml
Resolving tinyurl.com (tinyurl.com)... 104.20.218.42, 104.20.219.42,
Connecting to tinyurl.com (tinyurl.com)|104.20.218.42|:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://raw.githubusercontent.com/microservices-demo/microservices-...
--2017-11-02 16:54:27-- https://raw.githubusercontent.com/microservices-dem...
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.5...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101...
HTTP request sent, awaiting response... 200 OK
<output_omitted>
```

- Find the expected namespaces inside the file. It should be sock-shop. Also note the various settings. This file will deploy several containers which work together, providing a shopping website. As we work with other parameters you could revisit this file to see potential settings.

```
student@lfs458-node-1a0a:~$ less complete-demo.yaml
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: carts-db
  labels:
    name: carts-db
    namespace: sock-shop
spec:
  replicas: 1
<output_omitted>
```

- Create the namespace and verify it was made.

```
student@lfs458-node-1a0a:~$ kubectl create namespace sock-shop
namespace "sock-shop" created
```

```
student@lfs458-node-1a0a:~$ kubectl get namespace
NAME          STATUS    AGE
default       Active    35m
kube-public   Active    35m
kube-system   Active    35m
sock-shop     Active    5s
```

4. View the images the new application will deploy.

```
student@lfs458-node-1a0a:~$ grep image complete-demo.yaml
image: mongo
image: weaveworksdemos/carts:0.4.8
image: weaveworksdemos/catalogue-db:0.3.0
image: weaveworksdemos/catalogue:0.3.5
image: weaveworksdemos/front-end:0.3.12
image: mongo
image: weaveworksdemos/orders:0.4.7
image: weaveworksdemos/payment:0.4.3
image: weaveworksdemos/queue-master:0.3.1
image: rabbitmq:3.6.8
image: weaveworksdemos/shipping:0.4.8
image: weaveworksdemos/user-db:0.4.0
image: weaveworksdemos/user:0.4.4
```

5. Create the new shopping website using the YAML file. Use the namespace you recently created. Note that the deployments match the images we saw in the file.

```
student@lfs458-node-1a0a:~$ kubectl apply -n sock-shop -f complete-demo.yaml
deployment "carts-db" created
service "carts-db" created
deployment "carts" created
service "carts" created
<output_omitted>
```

6. Using the proper namespace will be important. This can be set on a per-command basis or as a shell parameter. Note the first command shows no pods. We must remember to pass the proper namespace. Some containers may not have fully downloaded or deployed by the time you run the command.

```
student@lfs458-node-1a0a:~$ kubectl get pods
No resources found.
```

```
student@lfs458-node-1a0a:~$ kubectl get pods -n sock-shop
```

NAME	READY	STATUS	RESTARTS	AGE
carts-511261774-c4jwv	1/1	Running	0	1m
carts-db-549516398-tw9zs	1/1	Running	0	1m
catalogue-4293036822-sp5kt	1/1	Running	0	1m
catalogue-db-1846494424-qzhvk	1/1	Running	0	1m
front-end-2337481689-6s65c	1/1	Running	0	1m
orders-208161811-1gc6k	1/1	Running	0	1m
orders-db-2069777334-4sp01	1/1	Running	0	1m
payment-3050936124-2cn2l	1/1	Running	0	1m
queue-master-2067646375-vzq77	1/1	Running	0	1m
rabbitmq-241640118-vk3m9	0/1	ContainerCreating	0	1m
shipping-3132821717-lm7kn	0/1	ContainerCreating	0	1m
user-1574605338-24xrb	0/1	ContainerCreating	0	1m
user-db-2947298815-lx9kp	1/1	Running	0	1m

7. Verify the shopping cart is exposing a web page. Use the public IP address of your AWS node (not the one derived from the prompt) to view the page. Note the external IP is not yet configured. Find the NodePort service. First try port 80 then try port 300001 as shown under the PORTS column.

```
student@lfs458-node-1a0a:~$ kubectl get svc -n sock-shop
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
carts	ClusterIP	10.100.154.148	<none>	80/TCP	1m
carts-db	ClusterIP	10.111.120.73	<none>	27017/TCP	1m
catalogue	ClusterIP	10.100.8.203	<none>	80/TCP	1m
catalogue-db	ClusterIP	10.111.94.74	<none>	3306/TCP	1m
front-end	NodePort	10.98.2.137	<none>	80:30001/TCP	1m
orders	ClusterIP	10.110.7.215	<none>	80/TCP	1m
orders-db	ClusterIP	10.106.19.121	<none>	27017/TCP	1m
payment	ClusterIP	10.111.28.218	<none>	80/TCP	1m
queue-master	ClusterIP	10.102.181.253	<none>	80/TCP	1m
rabbitmq	ClusterIP	10.107.134.121	<none>	5672/TCP	1m
shipping	ClusterIP	10.99.99.127	<none>	80/TCP	1m
user	ClusterIP	10.105.126.10	<none>	80/TCP	1m
user-db	ClusterIP	10.99.123.228	<none>	27017/TCP	1m

8. Check to see which node is running the containers. Note that the webserver is answering on a node which is not hosting the all the containers. First we check the master, then the second node. The containers should have to do with **kube proxy** services and **flannel**. The following is the **sudo docker ps** on both nodes.

```
student@lfs458-node-1a0a:~$ sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
d79fdb6f7343	quay.io/coreos/flannel@sha256:1b401bf0c30bada9a539389c3be652b58fe38463361edf488e6543c8761d4970	"/opt/bin/flanneld --"	7 hours ago	Up 7 hours		k8s_kube-flannel_kube-flannel-ds-h7tcv_kube-system_dd3be09c-a21f-11e7-acad-0a053b35c9b8_1
4206c95d0675	gcr.io/google_containers/kube-proxy-amd64@sha256:1509f2fc8a60501d604d21d983ed6f5d0ea40ccdd7cc6ba6c994389ef7db16d8	"/usr/local/bin/kube-proxy --"	7 hours ago	Up 7 hours		k8s_kube-proxy_kube-proxy-3xr5h_kube-system_dd3beec9-a21f-11e7-acad-0a053b35c9b8_0

```
student@lfs458-node-2b2b:~$ sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
b5a446a600e7	weaveworksdemos/user@sha256:6740c0a9a501c4d78e6f4e7ded2a7caa351b417df99	"/user -port=80"	3 minutes ago	Up 3 minutes		k8s_user_user-1574605338-24xrb_
sock-shop_1cde4173-a224-11e7-acad-0a053b35c9b8_0	3f97954048db	rabbitmq@sha256:a9f4923559bbcd00b93b02e61615aef5eb6f1d1c98db51053bab0fa6b680db26	"docker-entrypoint.sh"	3 minutes ago	Up 3 minutes	k8s_rabbitmq_rabbitmq-241640118-vk3m9_sock-shop_1c6e4f9c-a224-11e7-acad-0a053b35c9b8_0

<output_omitted>

9. Now we will shut down the shopping application. This can be done a few different ways. Begin by getting a listing of resources in all namespaces. There should be about 14 deployments.

```
student@lfs458-node-1a0a:~$ kubectl get deployment --all-namespaces
```

NAMESPACE	NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
kube-system	kube-dns	1	1	1	1	3h
low-usage-limit	limited-hog	1	1	1	1	33m
sock-shop	carts	1	1	1	1	6h
sock-shop	carts-db	1	1	1	1	6h
sock-shop	catalogue	1	1	1	1	6h

<output_omitted>

10. Use the terminal on the second node to get a count of the current docker containers. It should be something like 28, plus a line for status counted by **wc**. The main system should have something like 18 running, plus a line of status.

```
student@lfs458-node-2b2b:~$ sudo docker ps | wc -l
29

student@lfs458-node-1a0a:~$ sudo docker ps | wc -l
19
```

11. Delete some of the resources via deployments.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop delete deployment \
    catalogue catalogue-db front-end orders
deployment "catalogue" deleted
deployment "catalogue-db" deleted
```

12. Get a list of the pods that are running.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop get pod
NAME                                READY    STATUS    RESTARTS   AGE
carts-db-549516398-tw9zs            1/1     Running   0           6h
catalogue-4293036822-sp5kt          1/1     Running   0           6h
<output_omitted>
```

13. Delete a few resources using the pod name.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop delete pod \
    catalogue-4293036822-sp5kt catalogue-db-1846494424-qzhvk \
    front-end-2337481689-6s65c orders-208161811-1gc6k \
    orders-db-2069777334-4sp01
pod "catalogue-4293036822-sp5kt" deleted
pod "catalogue-db-1846494424-qzhvk" deleted
<output_omitted>
```

14. Check the status of the pods. There should be some pods running for only a few seconds. These will have the same name-stub as the Pods you recently deleted. The Deployment controller noticed expected number of Pods was not proper, so created new Pods until the current state matches the Pod manifest.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop get pod
NAME                                READY    STATUS    RESTARTS   AGE
catalogue-4293036822-mtz8m          1/1     Running   0           2s
catalogue-db-1846494424-16n2p        1/1     Running   0           22s
front-end-2337481689-6s65c          1/1     Terminating 0           6h
front-end-2337481689-80gwt          1/1     Running   0           22s
```

15. Delete the rest of the deployments. When no resources are found, examine the output of the `docker ps` command. None of the `sock-shop` containers should be found.

```
student@lfs458-node-1a0a:~$ kubectl get deployment --all-namespaces
NAMESPACE    NAME           DESIRED    CURRENT    UP-TO-DATE    AVAILABLE    AGE
kube-system   kube-dns       1          1          1             1            4h
low-usage-limit limited-hog    1          1          1             1            1h
sock-shop     carts          1          1          1             1            54m
sock-shop     carts-db       1          1          1             1            54m
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl delete deploy -n sock-shop carts carts-db....
<output_omitted>
```


Chapter 5

APIs and Access



5.1 Labs

Exercise 5.1: Configuring TLS Access

Overview

Using the Kubernetes API, **kubectl** makes API calls for you. With the appropriate TLS keys you could run **curl** as well use a **golang** client. Calls to the `kube-apiserver` get or set a `PodSpec`, or desired state. If the request represents a new state the **Kubernetes Control Plane** will update the cluster until the current state matches the specified state. Some end states may require multiple requests. For example, to delete a `ReplicaSet`, you would first set the number of replicas to zero, then delete the `ReplicaSet`.

An API request must pass information as JSON. **kubectl** converts `.yaml` to JSON when making an API request on your behalf. The API request has many settings, but must include `apiVersion`, `kind` and `metadata`, and `spec` settings to declare what kind of container to deploy. The `spec` fields depend on the object being created.

We will begin by configuring remote access to the `kube-apiserver` then explore more of the API.

Configuring TLS Access

1. Begin by reviewing the **kubectl** configuration file. We will use the three certificates and the API server address.

```
student@lfs458-node-1a0a:~$ less ~/.kube/config
<output_omitted>
```

2. We will set the certificates as variables. You may want to double-check each parameter as you set it. Begin with setting the `client-certificate-data` key.

```
student@lfs458-node-1a0a:~$ export client=$(grep client-cert ~/.kube/config |cut -d" " -f 6)

student@lfs458-node-1a0a:~$ echo $client
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUM4akNDQWRxZ0F3SUJ
```

```
BZ01JRy9wbC9rWEpNdm3RFFZSktvWklodmNOQVFFTEJRQXdGVEVUTUJFRO
ExVUUKQXhNS2EzVmlaWEp1WlhSbGN6QWVGdzB4TnpFeU1UTXh0elEyTXpKY
UZ3MHhPREV5TVRNeE56UTJNelJhTURReApGekFWQmdOVkJBb1REbk41YzNS
<output_omitted>
```

3. Almost the same command, but this time collect the `client-key-data` as the `key` variable.

```
student@lfs458-node-1a0a:~$ export key=$(grep client-key-data ~/.kube/config |cut -d " " -f 6)

student@lfs458-node-1a0a:~$ echo $key
<output_omitted>
```

4. Finally set the `auth` variable with the `certificate-authority-data` key.

```
student@lfs458-node-1a0a:~$ export auth=$(grep certificate-authority-data ~/.kube/config |cut -d " " -f 6)

student@lfs458-node-1a0a:~$ echo $auth
<output_omitted>
```

5. Now encode the keys for use with **curl**.

```
student@lfs458-node-1a0a:~$ echo $client | base64 -d - > ./client.pem

student@lfs458-node-1a0a:~$ echo $key | base64 -d - > ./client-key.pem

student@lfs458-node-1a0a:~$ echo $auth | base64 -d - > ./ca.pem
```

6. Pull the API server URL from the config file.

```
student@lfs458-node-1a0a:~$ kubectl config view |grep server
server: https://10.128.0.3:6443
```

7. Use **curl** command and the encoded keys to connect to the API server.

```
student@lfs458-node-1a0a:~$ curl --cert ./client.pem \
--key ./client-key.pem \
--cacert ./ca.pem \
https://10.128.0.3:6443/api/v1/pods

{
  "kind": "PodList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/pods",
    "resourceVersion": "239414"
  },
  <output_omitted>
```

8. If the previous command was successful, create a JSON file to create a new pod.

```
student@lfs458-node-1a0a:~$ vim curlpod.json
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    "namespace": "default",
    "labels": {
      "name": "examplepod"
    }
  },
  "spec": {
    "containers": [{
      "name": "nginx",
```

```

        "image": "nginx",
        "ports": [{"containerPort": 80}]
    }]
}
}

```

- The previous **curl** command can be used to build a XPOST API call. There will be a lot of output, including the scheduler and taints involved. Read through the output. In the last few lines the phase will probably show Pending, as it's near the beginning of the creation process.

```

student@lfs458-node-1a0a:~$ curl --cert ./client.pem \
--key ./client-key.pem --cacert ./ca.pem \
https://10.128.0.3:6443/api/v1/namespaces/default/pods \
-XPOST -H'Content-Type: application/json' \
-d@curlpod.json

{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",

```

- Verify the new pod exists.

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
curlpod       1/1     Running   0           15s

```

Exercise 5.2: Explore API Calls

- One way to view what a command does on your behalf is to use **strace**. In this case, we will look for the current endpoints, or targets of our API calls.

```

student@lfs458-node-1a0a:~$ kubectl get endpoints
NAME          ENDPOINTS          AGE
kubernetes    10.128.0.3:6443    3h

```

- Run this command again, preceded by **strace**. You will get a lot of output. Near the end you will note several **openat** functions to a local directory, `/home/student/.kube/cache/discovery/10.128.0.3_6443`

```

student@lfs458-node-1a0a:~$ strace kubectl get endpoints
execve("/usr/bin/kubectl", ["kubectl", "get", "endpoints"], [/*...
....
openat(AT_FDCWD, "/home/student/.kube/cache/discovery/10.128.0.3_6443..
<output_omitted>

```

- Change to the parent directory and explore. Your endpoint IP will be different, so replace the following with one suited to your system.

```

student@lfs458-node-1a0a:~$ cd /home/student/.kube/cache/discovery/

student@lfs458-node-1a0a:~/kube/cache/discovery$ ls
10.128.0.3_6443

student@lfs458-node-1a0a:~/kube/cache/discovery$ cd 10.128.0.3_6443/

```

- View the contents. You will find there are directories with various configuration information for kubernetes.

```

student@lfs458-node-1a0a:~/kube/cache/discovery/10.128.0.3_6443$ ls
apiextensions.k8s.io  batch                servergroups.json
apiregistration.k8s.io  certificates.k8s.io  settings.k8s.io
apps                  extensions            storage.k8s.io

```

```
authentication.k8s.io    networking.k8s.io        v1
authorization.k8s.io     policy
autoscaling              rbac.authorization.k8s.io
```

5. Use the find command to list out the subfiles. The prompt has been modified to look better on this page.

```
student@lfs458-node-1a0a:./10.128.0.3_6443$ find .
.
./authorization.k8s.io
./authorization.k8s.io/v1beta1
./authorization.k8s.io/v1beta1/serverresources.json
./authorization.k8s.io/v1
./authorization.k8s.io/v1/serverresources.json
./autoscaling
./autoscaling/v1
./autoscaling/v1/serverresources.json
<output_omitted>
```

6. View the objects available in version 1 of the API. For each object, or kind:, you can view the verbs or actions for that object, such as create seen in the following example. Note the prompt has been truncated for the command to fit on one line.

```
student@lfs458-node-1a0a:.$ python -m json.tool v1/serverresources.json
{
  "apiVersion": "v1",
  "groupVersion": "v1",
  "kind": "APIResourceList",
  "resources": [
    {
      "kind": "Binding",
      "name": "bindings",
      "namespaced": true,
      "singularName": "",
      "verbs": [
        "create"
      ]
    },
    <output_omitted>
```

7. Some of the objects have shortNames, which makes using them on the command line much easier. Locate the shortName for endpoints.

```
student@lfs458-node-1a0a:.$ python -m json.tool v1/serverresources.json | less
.
{
  "kind": "Endpoints",
  "name": "endpoints",
  "namespaced": true,
  "shortNames": [
    "ep"
  ],
  "singularName": "",
  "verbs": [
    "create",
    "delete",
  ]
}
```

8. Use the shortName to view the endpoints. It should match the output from the previous command.

```
student@lfs458-node-1a0a:.$ kubectl get ep
NAME          ENDPOINTS          AGE
kubernetes    10.128.0.3:6443    3h
```

9. We can see there are 37 objects in version 1 file.

```
student@lfs458-node-1a0a:.$ python -m json.tool v1/serverresources.json | grep kind
    "kind": "APIResourceList",
    "kind": "Binding",
    "kind": "ComponentStatus",
    "kind": "ConfigMap",
    "kind": "Endpoints",
    "kind": "Event",
<output_omitted>
```

10. Looking at another file we find seven more.

```
student@lfs458-node-1a0a:$ python -m json.tool \
    apps/v1beta1/serverresources.json | grep kind
    "kind": "APIResourceList",
    "kind": "ControllerRevision",
    "kind": "Deployment",
<output_omitted>
```

11. Delete the curlpod to recoup system resources.

```
student@lfs458-node-1a0a:$ kubectl delete po curlpod
pod "curlpod" deleted
```

12. Take a look around the other files as time permits.

Chapter 6

Managing State With Deployments



6.1 Labs

Exercise 6.1: Working with ReplicaSets

Overview

Understanding and managing the state of containers is a core Kubernetes task. In this lab we will first explore the API objects used to manage groups of containers. The objects available have changed as Kubernetes has matured, so the Kubernetes version in use will determine which are available. Our first object will be a `ReplicaSet`, which does not include newer management features found with `Deployments`. A `Deployment` will which manage `ReplicaSets` for you. We will also work with another object called a `DaemonSet` which ensures a container is running on newly added node.

Then we will update the software in a container, view the revision history, and roll-back to a previous version.

Working with ReplicaSets

A `ReplicaSet` is a next-generation of a `Replication Controller`, which differs only in the selectors supported. The only reason to use a `ReplicaSet` anymore is if you have no need for updating container software or require update orchestration which won't work with the typical process.

1. View any current `ReplicaSets`. If you deleted resources at the end of a previous lab, you should have none show from the command.

```
student@lfs458-node-1a0a:~$ kubectl get rs
No resources found.
```

2. Create a YAML file for a simple `ReplicaSet`. The `apiVersion` setting depends on the version of Kubernetes you are using. Versions 1.8 and beyond will use `apps/v1beta1`, then `apps/v1beta2` and then probably `apps/v1`.

```
student@lfs458-node-1a0a:~$ vim rs.yaml
```

```

apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: rs-one
spec:
  replicas: 2
  template:
    metadata:
      labels:
        system: ReplicaOne
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80

```

3. Create the ReplicaSet:

```

student@lfs458-node-1a0a:~$ kubectl create -f rs.yaml
replicaset "rs-one" created

```

4. View the newly created ReplicaSet:

```

student@lfs458-node-1a0a:~$ kubectl describe rs/rs-one
Name:                rs-one
Namespace:           default
Selector:             system=ReplicaOne
Labels:              system=ReplicaOne
Annotations:         <none>
Replicas:            2 current / 2 desired
Pods Status:         2 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:            system=ReplicaOne
  Containers:
    nginx:
      Image:          nginx:1.7.9
      Port:           80/TCP
      Environment:    <none>
      Mounts:         <none>
      Volumes:        <none>
  Events:            <none>

```

5. View the Pods created with the ReplicaSet. From the yaml file created there should be two Pods.

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME                READY    STATUS    RESTARTS   AGE
rs-one-2p9x4        1/1     Running   0          5m
rs-one-3c6pb        1/1     Running   0          5m

```

6. Now we will delete the ReplicaSet, but not the Pods it controls.

```

student@lfs458-node-1a0a:~$ kubectl delete rs/rs-one --cascade=false
replicaset "rs-one" deleted

```

View the ReplicaSet and Pods again:

7. student@lfs458-node-1a0a:~\$ kubectl describe rs/rs-one

Error from server (NotFound): replicaset.extensions "rs-one" not found

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME                READY    STATUS    RESTARTS   AGE
rs-one-2p9x4        1/1     Running   0          7m
rs-one-3c6pb        1/1     Running   0          7m

```


8. Create the ReplicaSet again. As long as we do not change the selector field, the new ReplicaSet should take ownership. Pod software versions cannot be updated this way.

```
student@lfs458-node-1a0a:~$ kubectl create -f rs.yaml
replicaset "rs-one" created
```

9. View the age of the ReplicaSet and then the Pods within:

```
student@lfs458-node-1a0a:~$ kubectl get rs
NAME          DESIRED  CURRENT  READY  AGE
rs-one        2        2        2      46s

student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS  AGE
rs-one-2p9x4  1/1     Running   0         8m
rs-one-3c6pb  1/1     Running   0         8m
```

10. We will now isolate a Pod from its ReplicaSet. Begin by editing the label of a Pod. We will change the system: parameter to be IsolatedPod.

```
student@lfs458-node-1a0a:~$ kubectl edit po rs-one-3c6pb

....
  labels:
    system: IsolatedPod
  name: rs-one-3c6pb
....
```

11. View the number of pods within the ReplicaSet. You should see two running.

```
student@lfs458-node-1a0a:~$ kubectl get rs
NAME          DESIRED  CURRENT  READY  AGE
rs-one        2        2        2      4m
```

12. New view the pods. You should note that there are three, with one being newer than others. The ReplicaSet made sure to keep two replicas, replacing the Pod which was isolated.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY   STATUS    RESTARTS  AGE
rs-one-2p9x4  1/1     Running   0         10m
rs-one-3c6pb  1/1     Running   0         10m
rs-one-dq5xd  1/1     Running   0         30s
```

13. Delete the ReplicaSet, then view any remaining Pods.

```
student@lfs458-node-1a0a:~$ kubectl delete rs/rs-one
replicaset "rs-one" deleted

student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY   STATUS    RESTARTS  AGE
rs-one-3c6pb  1/1     Running   0         14m
rs-one-dq5xd  0/1     Terminating 0         4m
```

14. In the above example the Pods had not finished termination. Wait for a bit and check again. There should be no ReplicaSets, but one Pod.

```
student@lfs458-node-1a0a:~$ kubectl get rs
No resources found.

student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY   STATUS    RESTARTS  AGE
rs-one-3c6pb  1/1     Running   0         16m
```

15. Delete the remaining Pod.

```
student@lfs458-node-1a0a:~$ kubectl delete po/rs-one-3c6pb
pod "rs-one-3c6pb" deleted
```

Exercise 6.2: Working with DaemonSets

A DaemonSet is a more flexible version of a Deployment which we have been working with in the rest of the labs. The DaemonSet ensures that as nodes are added to a cluster Pods will be created on them. A Deployment would only ensure a particular number of Pods are created in general. A node be should also be removed from a cluster; the DaemonSet would ensure the Pods are garbage collected before removal.

This extra step of automation can be useful for using with products like **ceph** where storage is often added or removed. They allow for complex deployments when used with declared resources like memory, CPU or volumes. While similar to **Replication Controllers** they are better used to allocate resources to particular systems.

1. We begin by creating a yaml file. In this case the kind would be set to DaemonSet. For ease of use we will copy the previously created `rs.yaml` file and make a couple edits. Remove the Replicas: 2 line. Be aware in future API versions, starting with v1.8, you may need to declare a Pod selector.

```
student@lfs458-node-1a0a:~$ cp rs.yaml ds.yaml

student@lfs458-node-1a0a:~$ vim ds.yaml
...
kind: DaemonSet
...
  name: ds-one
...
  system: DaemonSetOne
```

2. Create and verify the newly formed DaemonSet. There should be one Pod per node in the cluster.

```
student@lfs458-node-1a0a:~$ kubectl create -f ds.yaml
daemonset "ds-one" created

student@lfs458-node-1a0a:~$ kubectl get ds
NAME      DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE-SELECTOR   AGE
ds-one    2         2         2       2            2          <none>          1m

student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY   STATUS    RESTARTS   AGE
ds-one-b1dcv        1/1     Running   0          2m
ds-one-z31r4        1/1     Running   0          2m
```

3. Verify the image running inside the Pods. We will use this information in the next section.

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-one-b1dcv | grep Image:
Image:          nginx:1.7.9
```

Exercise 6.3: Rolling Updates and Rollbacks

One of the advantages of micro-services is the ability to replace and upgrade a container while continuing to respond to client requests. We will use the default `OnDelete` setting that upgrade a container when the predasor is deleted, then the use the `RollingUpdate` feature as well.

1. Begin by viewing the current updateStrategy setting for the DaemonSet created in the previous section.

```
student@lfs458-node-1a0a:~$ kubectl get ds/ds-one -o yaml \
| grep -A 1 Strategy
updateStrategy:
  type: OnDelete
```

2. Edit `ds.yaml`. Comment out the current image line and create a new line which uses version **1.8.1-alpine**.

```
student@lfs458-node-1a0a:~$ vim ds.yaml
....
- name: nginx
#   image: nginx:1.7.9
  image: nginx:1.8.1-alpine
  ports:
....
```

3. Use `kubectl replace` to update the configuration file in use for the DaemonSet.

```
student@lfs458-node-1a0a:~$ kubectl replace -f ds.yaml
daemonset "ds-one" replaced
```

4. Verify that the `Image:` parameter for the Pod checked in the previous section is unchanged.

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-one-b1dcv |grep Image:
Image:          nginx:1.7.9
```

5. Delete the Pod. Wait until the replacement Pod is running and check the version.

```
student@lfs458-node-1a0a:~$ kubectl delete po/ds-one-b1dcv
pod "ds-one-b1dcv" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY   STATUS    RESTARTS   AGE
ds-one-xc86w        1/1     Running   0           19s
ds-one-z31r4        1/1     Running   0           17m
```

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-one-xc86w |grep Image:
Image:          nginx:1.8.1-alpine
```

6. View the image running on the older Pod. It should still show version 1.7.9.

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-one-z31r4 |grep Image:
Image:          nginx:1.7.9
```

7. View the history of changes for the DaemonSet. You should see two revisions listed.

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds/ds-one
daemonsets "ds-one"
REVISION    CHANGE-CAUSE
1           <none>
2           <none>
```

8. View the settings for the various versions of the DaemonSet. The `Image:` line should be the only difference between the two outputs.

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds/ds-one --revision=1
daemonsets "ds-one" with revision #1
Pod Template:
  Labels:          system=DaemonSetOne
  Containers:
    nginx:
      Image:        nginx:1.7.9
      Port:         80/TCP
      Environment:  <none>
      Mounts:       <none>
      Volumes:      <none>

student@lfs458-node-1a0a:~$ kubectl rollout history ds/ds-one --revision=2
....
Image:          nginx:1.8.1-alpine
.....
```

9. Use `kubectl rollout undo ds/ds-one --to-revision=1` to change the DaemonSet back to an earlier version. As we are still using the `OnDelete` strategy there should be no change to the Pods.

```
student@lfs458-node-1a0a:~$ kubectl rollout undo ds/ds-one --to-revision=1
daemonset "ds-one" rolled back
```

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-one-xc86w |grep Image:
Image:          nginx:1.8.1-alpine
```

10. Delete the Pod, wait for the replacement to spawn then check the image version again.

```
student@lfs458-node-1a0a:~$ kubectl delete po/ds-one-xc86w
pod "ds-one-xc86w" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY   STATUS    RESTARTS   AGE
ds-one-qc72k        1/1     Running   0           10s
ds-one-xc86w        0/1     Terminating   0           12m
ds-one-z31r4        1/1     Running   0           28m
```

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-one-qc72k |grep Image:
Image:          nginx:1.7.9
```

11. View the details of the DaemonSet. The Image should be v1.7.9 in the output.

```
student@lfs458-node-1a0a:~$ kubectl describe ds |grep Image:
Image:          nginx:1.7.9
```

12. View the current configuration for the DaemonSet in YAML output. Look for the update strategy near the end of the output.

```
student@lfs458-node-1a0a:~$ kubectl get ds/ds-one -o yaml
apiVersion: extensions/v1beta1
kind: DaemonSet
.....
  terminationGracePeriodSeconds: 30
  templateGeneration: 3
  updateStrategy:
    type: OnDelete
status:
  currentNumberScheduled: 2
.....
```

13. Create a new DaemonSet, this time setting the update policy to `RollingUpdate`. Begin by generating a new config file.

```
student@lfs458-node-1a0a:~$ kubectl get ds/ds-one -o yaml > ds2.yaml
```

14. Edit the file. Change the following values and delete all the status lines at the bottom of the file. When finished the `type: RollingUpdate` line will probably be the last line in the file.

```
student@lfs458-node-1a0a:~$ vim ds2.yaml
....
  name: ds-two
....
  type: RollingUpdate
```

15. Create the new DaemonSet and verify the **nginx** version in the new pods.

```
student@lfs458-node-1a0a:~$ kubectl create -f ds2.yaml
daemonset "ds-two" created
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY    STATUS    RESTARTS   AGE
ds-one-qc72k        1/1      Running   0           28m
ds-one-z31r4        1/1      Running   0           57m
ds-two-10khc        1/1      Running   0           5m
ds-two-kzp9g        1/1      Running   0           5m
```

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-two-10khc |grep Image:
Image:                nginx:1.7.9
```

16. Edit the configuration file and set the image to a newer version such as 1.8.1-alpine.

```
student@lfs458-node-1a0a:~$ kubectl edit ds/ds-two
....
- image: nginx:1.8.1-alpine
.....
```

17. View the age of the DaemonSets. It should be around ten minutes old, depending on how fast you type.

```
student@lfs458-node-1a0a:~$ kubectl get ds/ds-two
NAME      DESIRED   CURRENT   READY     UP-TO-DATE   AVAILABLE   NODE-SELECTOR   AGE
ds-two    2         2         2         2            2           <none>          10m
```

18. Now view the age of the Pods. Two should be much younger than the DaemonSet.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY    STATUS    RESTARTS   AGE
ds-one-qc72k        1/1      Running   0           36m
ds-one-z31r4        1/1      Running   0           1h
ds-two-2p8vz        1/1      Running   0           34s
ds-two-8lx7k        1/1      Running   0           32s
```

19. Verify the Pods are using the new version of the software.

```
student@lfs458-node-1a0a:~$ kubectl describe po/ds-two-8lx7k |grep Image:
Image:                nginx:1.8.1-alpine
```

20. View the rollout status and the history of the DaemonSets.

```
student@lfs458-node-1a0a:~$ kubectl rollout status ds/ds-two
daemon set "ds-two" successfully rolled out
```

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds/ds-two
daemonsets "ds-two"
REVISION    CHANGE-CAUSE
1           <none>
2           <none>
```

21. View the changes in the update they should look the same as the previous history, but did not require the Pods to be deleted for the update to take place.

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds/ds-two --revision=2
...
Image:                nginx:1.8.1-alpine
```

22. Clean up the system by removing one of the DaemonSets. We will leave the other running.

```
student@lfs458-node-1a0a:~$ kubectl delete ds ds-two
daemonset "ds-two" deleted
```


Chapter 7

Volumes and Data



7.1 Labs

Exercise 7.1: Create a ConfigMap

Overview

Container files are ephemeral, which can be problematic for some applications. Should a container be restarted the files will be lost. In addition, we need a method to share files between containers inside a Pod.

A **Volume** is a directory accessible to containers in a Pod. Cloud providers offer volumes which persist further than the life of the Pod, such that AWS or GCE volumes could be pre-populated and offered to Pods, or transferred from one Pod to another. **Ceph** is also another popular solution for dynamic, persistent volumes.

Unlike current **Docker** volumes a Kubernetes volume has the lifetime of the Pod, not the containers within. You can also use different types of volumes in the same Pod simultaneously, but Volumes cannot mount in a nested fashion. Each must have their own mount point. Volumes are declared with `spec.volumes` and mount points with `spec.containers.volumeMounts` parameters. Each particular volume type, 24 currently, may have other restrictions. <https://kubernetes.io/docs/concepts/storage/volumes/#types-of-volumes>

We will also work with a **ConfigMap**, which is basically a set of key-value pairs. This data can be made available so that a Pod can read the data as environment variables or configuration data. A **ConfigMap** is similar to a **Secret**, except they are not base64 byte encoded arrays. They are stored as strings and can be read in serialized form.

Create a ConfigMap

There are three different ways a **ConfigMap** can ingest data, from a literal value, from a file or from a directory of files.

1. We will create a **ConfigMap** containing primary colors. We will create a series of files to ingest into the **ConfigMap**. First, we create a directory `primary` and populate it with four files. Then we create a file in our home directory with our favorite color.

```

student@lfs458-node-1a0a:~$ mkdir primary

student@lfs458-node-1a0a:~$ echo c > primary/cyan

student@lfs458-node-1a0a:~$ echo m > primary/magenta

student@lfs458-node-1a0a:~$ echo y > primary/yellow

student@lfs458-node-1a0a:~$ echo k > primary/black

student@lfs458-node-1a0a:~$ echo "known as key" >> primary/black

student@lfs458-node-1a0a:~$ echo blue > favorite

```

2. Now we will create the ConfigMap and populate it with the files we created as well as a literal value from the command line.

```

student@lfs458-node-1a0a:~$ kubectl create configmap colors \
    --from-literal=text=black \
    --from-file=./favorite \
    --from-file=./primary/
configmap "colors" created

```

3. View how the data is organized inside the cluster.

```

student@lfs458-node-1a0a:~$ kubectl get configmap colors
NAME      DATA      AGE
colors    6          30s

student@lfs458-node-1a0a:~$ kubectl get configmap colors -o yaml
apiVersion: v1
data:
  black: |
    k
    known as key
  cyan: |
    c
  favorite: |
    blue
  magenta: |
    m
  text: black
  yellow: |
    y
kind: ConfigMap
<output_omitted>

```

4. Now we can create a Pod to use the ConfigMap. In this case a particular parameter is being defined as an environment variable.

```

student@lfs458-node-1a0a:~$ vim simpleshell.yaml

apiVersion: v1
kind: Pod
metadata:
  name: shell-demo
spec:
  containers:
  - name: nginx
    image: nginx
    env:

```



```
- name: ilike
  valueFrom:
    configMapKeyRef:
      name: colors
      key: favorite
```

5. Create the Pod and view the environmental variable. After you view the parameter, exit out and delete the pod.

```
student@lfs458-node-1a0a:~$ kubectl create -f simpleshell.yaml
pod "shell-demo" created

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo \
    -- /bin/bash -c 'echo $ilike'
blue

student@lfs458-node-1a0a:~$ kubectl delete pod shell-demo
pod "shell-demo" deleted
```

6. All variables from a file can be included as environment variables as well. Comment out the previous `env:` stanza and add a slightly different `envFrom` to the file. Having new and old code at the same time can be helpful to see and understand the differences. Recreate the Pod, check all variables and delete the pod again. They can be found spread throughout the environment variable output.

```
student@lfs458-node-1a0a:~$ vim simpleshell.yaml
<output_omitted>
  image: nginx
#   env:
#   - name: ilike
#     valueFrom:
#       configMapKeyRef:
#         name: colors
#         key: favorite
  envFrom:
    - configMapRef:
      name: colors

student@lfs458-node-1a0a:~$ kubectl create -f simpleshell.yaml

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo \
    -- /bin/bash -c 'env'
HOSTNAME=shell-demo
NJS_VERSION=1.13.6.0.1.14-1~stretch
NGINX_VERSION=1.13.6-1~stretch
black=k
know as key

favorite=blue
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl delete pod shell-demo
pod "shell-demo" deleted
```

7. A ConfigMap can also be created from a YAML file. Create one with a few parameters to describe a car.

```
student@lfs458-node-1a0a:~$ vim car-map.yaml

apiVersion: v1
kind: ConfigMap
metadata:
  name: fast-car
  namespace: default
```

```
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby
```

8. Create the ConfigMap and verify the settings.

```
student@lfs458-node-1a0a:~$ kubectl create -f car-map.yaml
configmap "fast-car" created

student@lfs458-node-1a0a:~$ kubectl get configmap fast-car -o yaml
apiVersion: v1
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby
kind: ConfigMap
<output_omitted>
```

9. We will now make the ConfigMap available to a Pod as a mounted volume. You can again comment out the previous environmental settings and add the following new stanza. The containers: and volumes: entries are indented the same number of spaces.

```
student@lfs458-node-1a0a:~$ vim simpleshell.yaml
<output_omitted>
spec:
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
        - name: car-vol
          mountPath: /etc/cars
  volumes:
    - name: car-vol
      configMap:
        name: fast-car
```

10. Create the Pod again. Verify the volume exists and the contents of a file within. Due to the lack of a carriage return in the file your next prompt may be on the same line as the output, Shelby.

```
student@lfs458-node-1a0a:~$ kubectl create -f simpleshell.yaml
pod "shell-demo" created

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo -- \
  /bin/bash -c 'df -h'
Filesystem      Size  Used Avail Use% Mounted on
none            20G   1.8G   18G   10% /
tmpfs           3.9G    0   3.9G    0% /dev
tmpfs           3.9G    0   3.9G    0% /sys/fs/cgroup
/dev/xvda1      20G   1.8G   18G   10% /etc/cars
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo -- \
  /bin/bash -c 'cat /etc/cars/car.trim'
Shelby
```

11. Delete the Pod and ConfigMaps we were using.

```
student@lfs458-node-1a0a:~$ kubectl delete pods shell-demo
pod "shell-demo" deleted

student@lfs458-node-1a0a:~$ kubectl delete configmap fast-car colors
```

```
configmap "fast-car" deleted
configmap "colors" deleted
```

Exercise 7.2: Creating a Persistent NFS Volume (PV)

We will first deploy an NFS server. Once tested we will create a persistent NFS volume for containers to claim.

1. Install the software on your master node.

```
student@lfs458-node-1a0a:~$ sudo apt-get update && sudo \
    apt-get install -y nfs-kernel-server
<output_omitted>
```

2. Make and populate a directory to be shared. Also give it similar permissions to `/tmp/`

```
student@lfs458-node-1a0a:~$ sudo mkdir /opt/sfw

student@lfs458-node-1a0a:~$ sudo chmod 1777 /opt/sfw/

student@lfs458-node-1a0a:~$ sudo bash -c \
    'echo software > /opt/sfw/hello.txt'
```

3. Edit the NFS server file to share out the newly created directory. In this case we will share the directory with all. You can always **snoop** to see the inbound request in a later step and update the file to be more narrow.

```
student@lfs458-node-1a0a:~$ sudo vim /etc/exports
/opt/sfw/ *(rw,sync,no_root_squash,subtree_check)
```

4. Cause `/etc/exports` to be re-read:

```
student@lfs458-node-1a0a:~$ sudo exportfs -ra
```

5. Test by mounting the resource from your **second** node.

```
student@lfs458-node-2b2b:~$ sudo apt-get -y install nfs-common
<output_omitted>

student@lfs458-node-2b2b:~$ showmount -e lfs458-node-1a0a
Export list for lfs458-node-1a0a:
/opt/sfw *

student@lfs458-node-2b2b:~$ sudo mount 10.128.0.3:/opt/sfw /mnt

student@lfs458-node-2b2b:~$ ls -l /mnt
total 4
-rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt
```

6. Return to the master node and create a YAML file for the object with kind, PersistentVolume. Use the hostname of the master server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the resource will not start. Note that the accessModes do not currently affect actual access and are typically used as labels instead.

```
student@lfs458-node-1a0a:~$ vim PVol.yaml
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pvvol-1
spec:
  capacity:
```

```

    storage: 1Gi
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  nfs:
    path: /opt/sfw
    server: lfs458-node-1a0a    #<-- Edit to match master node
    readOnly: false

```

7. Create the persistent volume, then verify its creation.

```

student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml
persistentvolume "pvvol-1" created

student@lfs458-node-1a0a:~$ kubectl get pv
NAME          CAPACITY  ACCESSMODES  RECLAIMPOLICY  STATUS
CLAIM         STORAGECLASS  REASON      AGE
pvvol-1       1Gi         RWX          Retain          Available    4s

```

Exercise 7.3: Creating a Persistent Volume Claim (PVC)

Before Pods can take advantage of the new PV we need to create a **Persistent Volume Claim (PVC)**.

1. Begin by determining if any currently exist.

```

student@lfs458-node-1a0a:~$ kubectl get pvc
No resources found.

```

2. Create a YAML file for the new pvc.

```

student@lfs458-node-1a0a:~$ vim pvc.yaml

```

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-one
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 200Mi

```

3. Create and verify the new pvc is bound. Note that the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used.

```

student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml
persistentvolumeclaim "pvc-one" created

student@lfs458-node-1a0a:~$ kubectl get pvc
NAME      STATUS  VOLUME  CAPACITY  ACCESSMODES  STORAGECLASS  AGE
pvc-one   Bound   pvvol-1  1Gi       RWX           default/pvc-one  4s

```

4. Look at the status of the pv again, to determine if it is in use. It should show a status of Bound.

```

student@lfs458-node-1a0a:~$ kubectl get pv
NAME          CAPACITY  ACCESSMODES  RECLAIMPOLICY  STATUS  CLAIM                STORAGECLASS  REASON  AGE
pvvol-1       1Gi         RWX          Retain          Bound   default/pvc-one      default/pvc-one  5m

```

5. Create a new deployment to use the pvc. We will copy and edit an existing deployment yaml file. We will change the deployment name then add a volumeMounts section under containers and volumes section to the general spec. The name used must match in both places, whatever name you use. The claimName must match an existing pvc. As shown in the following example.

```
student@lfs458-node-1a0a:~$ cp first.yaml nfs-pod.yaml
```

```
student@lfs458-node-1a0a:~$ vim nfs-pod.yaml
```

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  generation: 1
  labels:
    run: nginx
  name: nginx-nfs
  namespace: default
  resourceVersion: "1411"
spec:
  replicas: 1
  selector:
    matchLabels:
      run: nginx
  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
    type: RollingUpdate
  template:
    metadata:
      creationTimestamp: null
      labels:
        run: nginx
    spec:
      containers:
      - image: nginx
        imagePullPolicy: Always
        name: nginx
        volumeMounts:
        - name: nfs-vol
          mountPath: /opt
        ports:
        - containerPort: 80
          protocol: TCP
        resources: {}
        terminationMessagePath: /dev/termination-log
        terminationMessagePolicy: File
      volumes:
      - name: nfs-vol
        persistentVolumeClaim:
          claimName: pvc-one
        dnsPolicy: ClusterFirst
        restartPolicy: Always
        schedulerName: default-scheduler
        securityContext: {}
        terminationGracePeriodSeconds: 30
```

6. Create the pod using the newly edited file.

```
student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml
```

- Look at the details of the pod.

```
student@lfs458-node-1a0a:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-nfs-1054709768-s8g28         1/1     Running   0           3m

student@lfs458-node-1a0a:~$ kubectl describe pod nginx-nfs-1054709768-s8g28
Name:                                nginx-nfs-1054709768-s8g28
Namespace:                          default
Node:                                lfs458-node-2b2b/10.128.0.5

<output_omitted>

Mounts:
  /opt from nfs-vol (rw)

<output_omitted>

Volumes:
  nfs-vol:
    Type:          PersistentVolumeClaim (a reference to a PersistentV...
    ClaimName:      pvc-one
    ReadOnly:       false
<output_omitted>
```

- View the status of the PVC. It should show as bound.

```
student@lfs458-node-1a0a:~$ kubectl get pvc
NAME      STATUS VOLUME  CAPACITY ACCESS MODES  STORAGECLASS  AGE
pvc-one   Bound  pvvol-1  1Gi      RWX              <empty>        2m
```

Exercise 7.4: Using a ResourceQuota to Limit PVC Count and Usage

The flexibility of cloud-based storage often requires limiting consumption among users. We will use the `ResourceQuota` object to both limit the total consumption as well as the number of persistent volume claims.

- Begin by deleting the deployment we had created to use NFS, the pv and the pvc.

```
student@lfs458-node-1a0a:~$ kubectl delete deploy nginx-nfs
deployment "nginx-nfs" deleted

student@lfs458-node-1a0a:~$ kubectl delete pvc pvc-one
persistentvolumeclaim "pvc-one" deleted

student@lfs458-node-1a0a:~$ kubectl delete pv pvvol-1
persistentvolume "pvvol-1" deleted
```

- Create a yaml file for the `ResourceQuota` object. Set the storage limit to ten claims with a total usage of 500Mi.

```
student@lfs458-node-1a0a:~$ vim storage-quota.yaml
```

- Create a new namespace called `small`. View the namespace information prior to the new quota. Either the long name with double dashes `--namespace` or the nickname `ns` work for the resource.

```
student@lfs458-node-1a0a:~$ kubectl create --namespace small
namespace "small" created

student@lfs458-node-1a0a:~$ kubectl describe ns small
Name:      small
Labels:    <none>
```

```
Annotations:      <none>
Status:           Active
```

No resource quota.

No resource limits.

4. Create a new pv and pvc in the small namespace.

```
student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml -n small
persistentvolume "pvvol-1" created
```

```
student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml -n small
persistentvolumeclaim "pvc-one" created
```

5. Create the new resource quota, placing this object into the low-usage-limit namespace.

```
student@lfs458-node-1a0a:~$ kubectl create -f storage-quota.yaml \
-n small
resourcequota "storagequota" created
```

6. Verify the small namespace has quotas. Compare the output to the same command above.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
Name:         small
Labels:       <none>
Annotations:  <none>
Status:       Active
```

Resource Quotas

Name:	storagequota	
Resource	Used	Hard
-----	---	---
persistentvolumeclaims	1	10
requests.storage	200Mi	500Mi

No resource limits.

7. Remove the namespace line from the `nfs-pod.yaml` file. Should be around line 11 or so. This will allow us to pass other namespaces on the command line.

```
student@lfs458-node-1a0a:~$ vim nfs-pod.yaml
```

8. Create the container again.

```
student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml \
-n small
deployment "nginx-nfs" created
```

9. Determine if the deployment has a running pod.

```
student@lfs458-node-1a0a:~$ kubectl get deploy --namespace=small
NAME          DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
nginx-nfs     1        1        1           0          43s
```

```
student@lfs458-node-1a0a:~$ kubectl describe deploy nginx-nfs \
-n small
<output_omitted>
```

10. Look to see if the pods are ready.

```
student@lfs458-node-1a0a:~$ kubectl get po --namespace=small
NAME                                READY    STATUS    RESTARTS   AGE
nginx-nfs-2854978848-g3khf         1/1      Running   0           37s
```

11. Ensure the Pod is running and is using the NFS mounted volume.

```
student@lfs458-node-1a0a:~$ kubectl describe po \
    nginx-nfs-2854978848-g3khf -n small
Name:                                nginx-nfs-2854978848-g3khf
Namespace:                          small
<output_omitted>

Mounts:
  /opt from nfs-vol (rw)
<output_omitted>
```

12. View the quota usage of the namespace

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>

Resource Quotas
Name:                                storagequota
Resource          Used    Hard
-----
persistentvolumeclaims    1      10
requests.storage          200Mi   500Mi

No resource limits.
```

13. Create a 300M file inside of the `/opt/sfw` directory on the host and view the quota usage again. Note that with NFS the size of the share is not counted against the deployment.

```
student@lfs458-node-1a0a:~$ sudo dd if=/dev/zero \
    of=/opt/sfw/bigfile bs=1M count=300
300+0 records in
300+0 records out
314572800 bytes (315 MB, 300 MiB) copied, 0.196794 s, 1.6 GB/s

student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
Resource Quotas
Name:                                storagequota
Resource          Used    Hard
-----
persistentvolumeclaims    1      10
requests.storage          200Mi   500Mi
<output_omitted>

student@lfs458-node-1a0a:~$ du -h /opt/
301M    /opt/sfw
41M     /opt/cni/bin
41M     /opt/cni
341M    /opt/
```

14. Now let us illustrate what happens when a deployment requests more than the quota. Begin by shutting down the existing deployment.

```
student@lfs458-node-1a0a:~$ kubectl get deploy -n small
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx-nfs     1         1         1             1           11m

student@lfs458-node-1a0a:~$ kubectl delete deploy nginx-nfs -n small
deployment "nginx-nfs" deleted
```


15. Once the Pod has shut down view the resource usage of the namespace again. Note the storage did not get cleaned up when the pod was shut down.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
Resource Quotas
Name:
Resource          storagequota
-----
persistentvolumeclaims    1      10
requests.storage          200Mi   500Mi
```

16. Remove the pvc then view the pv it was using. Note the RECLAIM POLICY and STATUS.

```
student@lfs458-node-1a0a:~$ kubectl get pvc -n small
NAME      STATUS    VOLUME    CAPACITY   ACCESSMODES  STORAGECLASS  AGE
pvc-one   Bound     pvvol-1   1Gi        RWX           small         19m

student@lfs458-node-1a0a:~$ kubectl delete pvc pvc-one -n small
persistentvolumeclaim "pvc-one" deleted

student@lfs458-node-1a0a:~$ kubectl get pv -n small
NAME      CAPACITY   ACCESSMODES  RECLAIMPOLICY  STATUS    CLAIM
STORAGECLASS  REASON    AGE
pvvol-1   1Gi    RWX    Retain   Released   small/pvc-one   44m
```

17. Dynamically provisioned storage uses the ReclaimPolicy of the StorageClass which could be Delete, Retain, or some types allow Recycle. Manually created persistent volumes default to Retain unless set otherwise at creation. The default storage policy is to retain the storage to allow recovery of any data. To change this begin by viewing the yaml output.

```
student@lfs458-node-1a0a:~$ kubectl get pv/pvvol-1 -o yaml
....
  path: /opt/sfw
  server: lfs458-node-1a0a
  persistentVolumeReclaimPolicy: Retain
status:
  phase: Released
```

18. Currently we will need to delete and re-create the object. Future development on a deleter plugin is planned. We will re-create the volume and allow it to use the Retain policy, then change it once running.

```
student@lfs458-node-1a0a:~$ kubectl delete pv/pvvol-1
persistentvolume "pvvol-1" deleted

student@lfs458-node-1a0a:~$ grep Retain PVol.yaml
  persistentVolumeReclaimPolicy: Retain

student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml
persistentvolume "pvvol-1" created
```

19. We will use kubectl patch to change the retention policy to Delete. The yaml output from before can be helpful in getting the correct syntax.

```
student@lfs458-node-1a0a:~$ kubectl patch pv pvvol-1 -p \
'{"spec":{"persistentVolumeReclaimPolicy":"Delete"}}'
persistentvolume "pvvol-1" patched

student@lfs458-node-1a0a:~$ kubectl get pv/pvvol-1
NAME      CAPACITY   ACCESSMODES  RECLAIMPOLICY  STATUS    CLAIM
STORAGECLASS  REASON    AGE
pvvol-1   1Gi    RWX    Delete   Available   2m
```

20. View the current quota settings.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
.
requests.storage          0          500Mi
```

21. Create the pvc again. Even with no pods running, note the resource usage.

```
student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml -n small
persistentvolumeclaim "pvc-one" created

student@lfs458-node-1a0a:~$ kubectl describe ns small
.
requests.storage          200Mi       500Mi
```

22. Remove the existing quota from the namespace.

```
student@lfs458-node-1a0a:~$ kubectl get resourcequota -n small
NAME          AGE
storagequota  12m

student@lfs458-node-1a0a:~$ kubectl delete resourcequota \
    storagequota -n small
resourcequota "storagequota" deleted
```

23. Edit the storagequota.yaml file and lower the capacity to 100Mi.

```
student@lfs458-node-1a0a:~$ vim storage-quota.yaml

.
requests.storage: "100Mi"
```

24. Create and verify the new storage quota. Note the hard limit has already been exceeded.

```
student@lfs458-node-1a0a:~$ kubectl create -f storage-quota.yaml -n small
resourcequota "storagequota" created

student@lfs458-node-1a0a:~$ kubectl describe ns small
.
persistentvolumeclaims      1          10
requests.storage            200Mi       100Mi

No resource limits.
```

25. Create the deployment again. View the deployment. Note there are no errors seen.

```
student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml \
    --namespace=small
deployment "nginx-nfs" created

student@lfs458-node-1a0a:~$ kubectl describe deploy/nginx-nfs -n small
Name:          nginx-nfs
Namespace:     small
<output_omitted>
```

26. Examine the pods to see if they are actually running.

```
student@lfs458-node-1a0a:~$ kubectl get po -n small
NAME                                READY    STATUS    RESTARTS   AGE
nginx-nfs-2854978848-vb6bh          1/1      Running   0           58s
```

27. As we were able to deploy more pods even with apparent hard quota set, let us test to see if the reclaim of storage takes place. Remove the deployment and the persistent volume claim.

```
student@lfs458-node-1a0a:~$ kubectl delete deploy nginx-nfs -n small
deployment "nginx-nfs" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl delete pvc/pvc-one -n small
persistentvolumeclaim "pvc-one" deleted
```

28. View if the persistent volume exists. You will see it attempted a removal, but failed. If you look closer you will find the error has to do with the lack of a deleter volume plugin for NFS. Other storage protocols have a plugin.

```
student@lfs458-node-1a0a:~$ kubectl get pv -n small
NAME          CAPACITY  ACCESSMODES  RECLAIMPOLICY  STATUS  CLAIM
STORAGECLASS  REASON    AGE
pvvol-1       1Gi      RWX          Delete         Failed  small/pvc-one  20m
```

29. Ensure the deployment, pvc and pv are all removed.

```
student@lfs458-node-1a0a:~$ kubectl delete pv/pvvol-1
persistentvolume "pvvol-1" deleted
```

30. Edit the persistent volume YAML file and change the persistentVolumeReclaimPolicy: to Recycle.

```
student@lfs458-node-1a0a:~$ vim PVol.yaml
....
persistentVolumeReclaimPolicy: Recycle
....
```

31. Add a LimitRange to the namespace and attempt to create the persistent volume and persistent volume claim again. We can use the LimitRange we used earlier.

```
student@lfs458-node-1a0a:~$ kubectl create -f low-resource-range.yaml\
-n small
limitrange "low-resource-range" created
```

32. View the settings for the namespace. Both quotas and resource limits should be seen.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
Resource Limits
Type      Resource  Min  Max  Default Request  Default Limit  ...
----      -
Container cpu      -    -    500m             1              -
Container memory   -    -    100Mi            500Mi          -
```

33. Create the persistent volume again. View the resource. Note the Reclaim Policy is Recycle.

```
student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml -n small
persistentvolume "pvvol-1" created
```

```
student@lfs458-node-1a0a:~$ kubectl get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  ...
pvvol-1       1Gi      RWX           Recycle         Available  ...
```

34. Attempt to create the persistent volume claim again. The quota only takes effect if there is also a resource limit in effect.

```
student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml -n small
Error from server (Forbidden): error when creating "pvc.yaml":
persistentvolumeclaims "pvc-one" is forbidden: exceeded quota:
storagequota, requested: requests.storage=200Mi, used:
requests.storage=0, limited: requests.storage=100Mi
```

35. Edit the resourcequota to increase the `requests.storage` to 500mi.

```
student@lfs458-node-1a0a:~$ kubectl edit resourcequota -n small
....
spec:
  hard:
    persistentvolumeclaims: "10"
    requests.storage: 500Mi
status:
  hard:
    persistentvolumeclaims: "10"
....
```

36. Create the pvc again. It should work this time. Then create the deployment again.

```
student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml -n small
persistentvolumeclaim "pvc-one" created

student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml -n small
deployment "nginx-nfs" created
```

37. View the namespace settings.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
```

38. Delete the deployment. View the status of the pv and pvc.

```
student@lfs458-node-1a0a:~$ kubectl delete deploy nginx-nfs -n small
deployment "nginx-nfs" deleted

student@lfs458-node-1a0a:~$ kubectl get pvc -n small
NAME      STATUS    VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE
pvc-one   Bound     pvvol-1   1Gi        RWX              small           7m

student@lfs458-node-1a0a:~$ kubectl get pv -n small
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM   STORA...
pvvol-1   1Gi        RWX            Recycle Bound     small   small/pvc-one   ...
```

39. Delete the pvc and check the status of the pv. It should show as Available.

```
student@lfs458-node-1a0a:~$ kubectl delete pvc pvc-one -n small
persistentvolumeclaim "pvc-one" deleted

student@lfs458-node-1a0a:~$ kubectl get pv -n small
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM   STORA...
pvvol-1   1Gi        RWX            Recycle          Available   ...
```

40. Remove the pv and any other resources created during this lab.

```
student@lfs458-node-1a0a:~$ kubectl delete pv pvvol-1
persistentvolume "pvvol-1" deleted
```

Chapter 8

Services



8.1 Labs

Exercise 8.1: Deploy A New Service

Overview

Services (also called **microservices**) are objects which declare a policy to access a logical set of Pods. They are typically assigned with `labels` to allow persistent access to a resource, when front or back end containers are terminated and replaced.

Native applications can use the `Endpoints` API for access. Non-native applications can use a Virtual IP-based bridge to access back end pods. `ServiceTypes` Type could be:

- **ClusterIP** default - exposes on a cluster-internal IP. Only reachable within cluster
- **NodePort** Exposes node IP at a static port. A ClusterIP is also automatically created.
- **LoadBalancer** Exposes service externally using cloud providers load balancer. NodePort and ClusterIP automatically created.
- **ExternalName** Maps service to contents of `externalName` using a CNAME record.

We use services as part of decoupling such that any agent or object can be replaced without interruption to access from client to back end application.

Deploy A New Service

1. Deploy two **nginx** servers using **kubectl** and a new `.yaml` file. We will use the `v1beta` version of the API. The kind should be `Deployment` and label it with `nginx`. Create two replicas and expose port 8080. What follows is a well documented file. There is no need to include the comments when you create the file.

```
student@lfs458-node-1a0a:~$ vim nginx-one.yaml
```

```

apiVersion: extensions/v1beta1
# Determines YAML versioned schema.
kind: Deployment
# Describes the resource defined in this file.
metadata:
  name: nginx-one
  labels:
    system: secondary
# Required string which defines object within namespace.
  namespace: accounting
# Existing namespace resource will be deployed into.
spec:
  replicas: 2
# How many Pods of following containers to deploy
  template:
    metadata:
      labels:
        app: nginx
# Some string meaningful to users, not cluster. Keys
# must be unique for each object. Allows for mapping
# to customer needs.
    spec:
      containers:
# Array of objects describing containerized application with a Pod.
# Referenced with shorthand spec.template.spec.containers
        - image: nginx:1.7.9
# The Docker image to deploy
          imagePullPolicy: Always
          name: nginx
# Unique name for each container, use local or Docker repo image
          ports:
            - containerPort: 8080
              protocol: TCP
# Optional resources this container may need to function.
          nodeSelector:
            system: secondOne
# One method of node affinity.

```

2. View the existing labels on the nodes in the cluster.

```
student@lfs458-node-1a0a:~$ kubectl get nodes --show-labels
```

3. Run the following command and look for the errors. Assuming there is no typo, you should have gotten an error about the accounting namespace.

```

student@lfs458-node-1a0a:~$ kubectl create -f nginx-one.yaml
Error from server (NotFound): error when creating
"nginx-one.yaml": namespaces "accounting" not found

```

4. Create the namespace and try to create the deployment again. There should be no errors this time.

```

student@lfs458-node-1a0a:~$ kubectl create ns accounting
namespace "accounting" created

student@lfs458-node-1a0a:~$ kubectl create -f nginx-one.yaml
deployment "nginx-one" created

```

5. View the status of the new nodes. Note they do not show a Running status.

```

student@lfs458-node-1a0a:~$ kubectl get pods -n accounting
NAME                                READY   STATUS    RESTARTS   AGE
nginx-one-74dd9d578d-fcpmv          0/1     Pending   0           4m
nginx-one-74dd9d578d-r2d67          0/1     Pending   0           4m

```

6. View the node each has been assigned to (or not) and the reason, which shows under events at the end of the output.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting describe pod
nginx-one-74dd9d578d-fcpmvlfs458-node-2b2b
Name:          nginx-one-74dd9d578d-fcpmv
Namespace:     accounting
Node:          <none>

<output_omitted>

Events:
  Type      Reason             Age          From          ....
  ----      -
Warning    FailedScheduling   3s (x25 over 6m) default-scheduler
No nodes are available that match all of the predicates:
MatchNodeSelector (2).
```

7. Label the secondary node. Verify the labels.

```
student@lfs458-node-1a0a:~$ kubectl label node lfs458-node-2b2b \
                        system=secondOne
node "lfs458-node-2b2b" labeled

student@lfs458-node-1a0a:~$ kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE      VERSION    LABELS
lfs458-node-1a0a    Ready    master   1d       v1.9.1    \
                    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/
hostname=lfs458-node-1a0a,node-role.kubernetes.io/master=
lfs458-node-2b2b    Ready    <none>   1d       v1.9.1    \
                    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/
hostname=lfs458-node-2b2b,system=secondOne
```

8. View the pods in the accounting namespace. They may still show as Pending. Depending on how long it has been since you attempted deployment the system may not have checked for the label. If the Pods show Pending after a minute delete one of the pods. They should both show as Running after as a deletion. A change in state will cause the Deployment controller to check the status of both Pods.

```
student@lfs458-node-1a0a:~$ kubectl get pods -n accounting
NAME                READY    STATUS    RESTARTS   AGE
nginx-one-74dd9d578d-fcpmv    1/1      Running    0           10m
nginx-one-74dd9d578d-sts5l    1/1      Running    0           3s
```

9. View Pods by the label we set in the YAML file. If you look back the Pods were given a label of app=nginx.

```
student@lfs458-node-1a0a:~$ kubectl get pods -l app=nginx --all-namespaces
NAMESPACE   NAME                READY    STATUS    RESTARTS   AGE
accounting  nginx-one-74dd9d578d-fcpmv    1/1      Running    0           20m
accounting  nginx-one-74dd9d578d-sts5l    1/1      Running    0           9m
```

10. Recall that we exposed port 8080 in the YAML file. Expose the new deployment.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment/nginx-one \
-n accounting
service "nginx-one" exposed
```

11. View the newly exposed endpoints. Note that port 8080 has been exposed on each Pod.

```
student@lfs458-node-9q6r:~$ kubectl get ep nginx-one -n accounting
NAME      ENDPOINTS                                AGE
nginx-one  10.244.0.21:8080,10.244.0.22:8080      11s
```

12. Attempt to access the Pod on port 8080, then on port 80. Even though we exposed port 8080 of the container the application within has not been configured to listen on this port. The **nginx** server will listen on port 80 by default. A `curl` command to that port should return the typical welcome page.

```
student@lfs458-node-1a0a:~$ curl 10.244.0.21:8080
curl: (7) Failed to connect to 10.244.0.21 port 8080: Connection refused

student@lfs458-node-1a0a:~$ curl 10.244.0.21:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

13. Delete the deployment. Edit the YAML file to expose port 80 and create the deployment again.

```
student@lfs458-node-1a0a:~$ kubectl delete deploy nginx-one -n accounting

student@lfs458-node-1a0a:~$ vim nginx-one.yaml

student@lfs458-node-1a0a:~$ kubectl create -f nginx-one.yaml
```

Exercise 8.2: Configure a NodePort

In a previous exercise we deployed a `LoadBalancer` which deployed a `ClusterIP` and `NodePort` automatically. In this exercise we will deploy a `NodePort`. While you can access a container from within the cluster, one can use a `NodePort` to NAT traffic from outside the cluster. One reason to deploy a `NodePort` instead, is that a `LoadBalancer` is also a load balancer resource from cloud providers like GKE and AWS.

1. In a previous step we were able to view the **nginx** page using the internal Pod IP address. Now expose the deployment using the `--type=NodePort`. We will also give it an easy to remember name and place it in the `accounting` namespace. We could pass the port as well, which could help with opening ports in the firewall.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment nginx-one \
  --type=NodePort --name=service-lab -n accounting
service "service-lab" exposed
```

2. View the details of the services in the `accounting` namespace. We are looking for the autogenerated port.

```
student@lfs458-node-1a0a:~$ kubectl describe services -n accounting
....
NodePort:                <unset>    32103/TCP
....
```

3. Locate the exterior facing IP address of the cluster. As we are using GCP nodes, which we access via a `FloatingIP`, we will first check the internal only public IP address. Look for the Kubernetes master URL.

```
student@lfs458-node-1a0a:~$ kubectl cluster-info
Kubernetes master is running at https://10.128.0.3:6443
KubeDNS is running at https://10.128.0.3:6443/api/v1/namespaces/kube-system/services/kube-dns/proxy
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

4. Test access to the **nginx** web server using the combination of master URL and `NodePort`.

```
student@lfs458-node-1a0a:~$ curl http://10.128.0.3:32103
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```


- Using the browser on your local system, use the public IP address you use to SSH into your node and the port. You should still see the **nginx** default page.

Exercise 8.3: Use Labels to Manage Resources

- Try to delete all Pods with the `app=nginx` label, in all namespaces. You should receive an error as this function must be narrowed to a particular namespace. Then delete using the appropriate namespace.

```
student@lfs458-node-1a0a:~$ kubectl delete pods -l app=nginx \
--all-namespaces
Error: unknown flag: --all-namespaces
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl delete pods -l app=nginx -n accounting
pod "nginx-one-74dd9d578d-fcpmv" deleted
pod "nginx-one-74dd9d578d-sts5l" deleted
```

- View the Pods again. New versions of the Pods should be running as the controller responsible for them continues.

```
student@lfs458-node-1a0a:~$ kubectl get pods -n accounting
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-one-74dd9d578d-ddt5r	1/1	Running	0	1m
nginx-one-74dd9d578d-hfzml	1/1	Running	0	1m

- We also gave a label to the deployment. View the deployment in the accounting namespace.

```
student@lfs458-node-1a0a:~$ kubectl get deploy -n accounting --show-labels
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	LABELS
nginx-one	2	2	2	2	27m	system=secondOne

- Delete the deployment using its label.

```
student@lfs458-node-1a0a:~$ kubectl delete deploy \
-l system=secondary -n accounting
deployment "nginx-one" deleted
```

- Remove the label from the secondary node. Note that the syntax is a minus sign directly after the key you want to remove, or `system` in this case.

```
student@lfs458-node-1a0a:~$ kubectl label node lfs458-node-2b2b system-
node "lfs458-node-2b2b" labeled.
```


Chapter 9

Ingress



9.1 Labs

There is no lab to complete for this chapter.

Chapter 10

API Objects



10.1 Labs

Exercise 10.1: RESTful API Access

Overview

We will continue to explore ways of accessing the control plane of our cluster. In the security chapter we will discuss there are several authentication methods, one of which is use of a Bearer token. We will work with one then deploy a local proxy server for application-level access to the Kubernetes API.

RESTful API Access

We will use the **curl** command to make API requests to the cluster, in an in-secure manner. Once we know the IP address and port, then the token we can retrieve cluster data in a RESTful manner. By default most of the information is restricted, but changes to authentication policy could allow more access.

1. First we need to know the IP and port of a node running a replica of the API server. The master system will typically have one running. Use **kubectl config view** to get overall cluster configuration, and find the server entry. This will give us both the IP and the port.

```
student@lfs458-node-1a0a:~$ kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
    server: https://10.128.0.3:6443
  name: kubernetes
<output_omitted>
```

2. Next we need to find the bearer token. This is part of a default token. Look at a list of tokens, first all on the cluster, then just those in the default namespace. There will be a secret for each of the controllers of the cluster.

- Pods can also make use of included certificates to use the API. The certificates are automatically made available to a pod under the `/var/run/secrets/kubernetes.io/serviceaccount/`. We will deploy a simple Pod and view the resources. If you view the `token` file you will find it is the same value we put into the `$token` variable.

```
student@lfs458-node-1a0a:~$ kubectl run -i -t busybox --image=busybox \
--restart=Never
# ls /var/run/secrets/kubernetes.io/serviceaccount/
ca.crt      namespace  token
/ # exit
```

Exercise 10.2: Using the Proxy

Another way to interact with the API is via a proxy. The proxy can be run from a node or from within a Pod through the use of a sidecar.

- Begin by starting the proxy. It will start in the foreground by default. There are several options you could pass. Begin by reviewing the help output.

```
student@lfs458-node-1a0a:~$ kubectl proxy -h
Creates a proxy server or application-level gateway between localhost
and the Kubernetes API Server. It also allows serving static content
over specified HTTP path. All incoming data enters through one port
and gets forwarded to the remote kubernetes API Server port, except
for the path matching the static content path.
```

Examples:

```
# To proxy all of the kubernetes api and nothing else, use:
```

```
$ kubectl proxy --api-prefix=/
<output_omitted>
```

- Start the proxy while setting the API prefix, and put it in the background.

```
student@lfs458-node-1a0a:~$ kubectl proxy --api-prefix=/ &
[1] 22500
Starting to serve on 127.0.0.1:8001
```

- Now use the same `curl` command, but point toward the IP and port shown by the proxy. The output should be the same as without the proxy.

```
student@lfs458-node-1a0a:~$ curl http://127.0.0.1:8001/api/
```

- Make an API call to retrieve the namespaces. The command did not work before due to permissions, but should work now as the proxy is making the request on your behalf.

```
student@lfs458-node-1a0a:~$ curl http://127.0.0.1:8001/api/v1/namespaces
{
  "kind": "NamespaceList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/namespaces",
    "resourceVersion": "86902"
  }
}
<output_omitted>
```

Exercise 10.3: Working with Cron Jobs

We will create a simple cron job to explore how to create them and view their execution. We will run a regular job and view both the job status as well as output. Note that the jobs are expected to be idempotent, so should not be used for tasks that require strict timings to run.

- Begin by creating a YAML for for the cron job. Set the time interval to be every minutes. Use the `busybox` container and pass it the `date` command.

```
student@lfs458-node-1a0a:~$ vim cron-job.yaml
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: date
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: dateperminute
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; sleep 30
          restartPolicy: OnFailure
```

2. Create the cron job.

```
student@lfs458-node-9q6r:~$ kubectl create -f cron-job.yaml
cronjob "date" created
```

3. Wait a few seconds for the job to ingested, then check that the job is active.

```
student@lfs458-node-9q6r:~$ kubectl get cronjob date
NAME      SCHEDULE      SUSPEND   ACTIVE   LAST SCHEDULE   AGE
date      */1 * * * *   False     1        Wed, 20 Dec 2017 15:56:00 +0000
```

4. The first attempt at the job may fail, depending on when in the time configured it was created. Also note that the output below shows that other jobs failed. Pass the **-watch** argument to view the jobs as they are created. Use **ctrl-c** to regain access to the shell.

```
student@lfs458-node-9q6r:~$ kubectl get jobs --watch
NAME          DESIRED   SUCCESSFUL   AGE
date-1513785360 1         1           19s
date-1513785420 1         0            0s
date-1513785420 1         1           33s
date-1513785480 1         0            0s
^C
```

5. Find Pods, including completed, associated with the date cronjob.

```
student@lfs458-node-9q6r:~$ kubectl get pods -a |grep date
date-1513786020-ws9j5    0/1      Completed   0      1m
date-1513786080-s7wfc    0/1      Completed   0      53s
```

6. Choose a particular pods and review its log. It should show the output of the data command when it executed.

```
student@lfs458-node-9q6r:~$ kubectl logs date-1513786080-s7wfc
Wed Dec 20 16:08:06 UTC 2017
```

7. When finished with the cronjob delete it.

```
student@lfs458-node-9q6r:~$ kubectl delete cronjob date
cronjob "date" deleted
```


Chapter 11

Scheduling



11.1 Labs

Exercise 11.1: Assign Pods Using Labels

Overview

While allowing the system to distribute Pods on your behalf is typically the best route, you may want to determine which nodes a Pod will use. For example you may have particular hardware requirements to meet for the workload. You may want to assign VIP Pods to new, faster hardware and everyone else to older hardware.

In this exercise we will use `labels` to schedule Pods to a particular node. Then we will explore `taints` to have more flexible deployment in a large environment.

Assign Pods Using Labels

1. Begin by getting a list of the nodes. They should be in the ready state and without added labels or taints.

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
lfs458-node-2b2b    Ready     <none>   2d       v1.9.1
lfs458-node-1a0a    Ready     master   2d       v1.9.1
```

2. View the current labels and taints for the nodes.

```
student@lfs458-node-1a0a:~$ kubectl describe nodes |grep -i label
Labels:                beta.kubernetes.io/arch=amd64
Labels:                beta.kubernetes.io/arch=amd64

student@lfs458-node-1a0a:~$ kubectl describe nodes |grep -i taint
Taints:                <none>
Taints:                <none>
```

3. Verify there are no deployments running. If there are, delete them. Get a count of how many container are running on both the master and secondary nodes. There are about 17 containers running on the master in the following example, and five running on the secondary. There are status lines which increase the **wc** count. You may have more or less, depending on previous labs and cleaning up of resources.

```
student@lfs458-node-1a0a:~$ kubectl get deployments --all-namespaces
NAMESPACE    NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
kube-system   kube-dns   1         1         1             1           1d
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
19
```

```
student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
7
```

4. For the purpose of the exercise we will assign the master node to be VIP hardware and the secondary node to be for others.

```
student@lfs458-node-1a0a:~$ kubectl label nodes lfs458-node-1a0a status=vip
node "lfs458-node-1a0a" labeled
```

```
student@lfs458-node-1a0a:~$ kubectl label nodes lfs458-node-2b2b status=other
node "lfs458-node-2b2b" labeled
```

5. Verify your settings. You will also find there are some built in labels such as hostname, os and architecture type. The output below appears on multiple line for readability.

```
student@lfs458-node-1a0a:~$ kubectl get nodes --show-labels
NAME                STATUS    AGE      VERSION   LABELS
lfs458-node-2b2b    Ready     2d       v1.9.1    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=lfs458-node-2b2b,status=other
lfs458-node-1a0a    Ready     2d       v1.9.1    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=lfs458-node-1a0a,node-role.kubernetes.io/master=,status=vip
```

6. Create **vip.yaml** to spawn four busybox containers which sleep the whole time. Include the **nodeSelector** entry.

```
student@lfs458-node-1a0a:~$ vim vip.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: vip
spec:
  containers:
  - name: vip1
    image: busybox
    args:
    - sleep
    - "1000000"
  - name: vip2
    image: busybox
    args:
    - sleep
    - "1000000"
  - name: vip3
    image: busybox
    args:
```

```

- sleep
- "1000000"
- name: vip4
  image: busybox
  args:
  - sleep
  - "1000000"
nodeSelector:
  status: vip

```

7. Deploy the new pod. Verify the containers have been created on the master node. It may take a few seconds for all the containers to spawn. Check both the master and the secondary nodes.

```

student@lfs458-node-1a0a:~$ kubectl create -f vip.yaml
pod "vip" created

```

```

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
24

```

```

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
7

```

8. Delete the pod then edit the file, commenting out the nodeSelector lines. It may take a while for the containers to fully terminate.

```

student@lfs458-node-1a0a:~$ kubectl delete pod vip
pod "vip" deleted

```

```

student@lfs458-node-1a0a:~$ vim vip.yaml
....
# nodeSelector:
#   status: vip

```

9. Create the pod again. Containers should now be spawning on both nodes.

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
vip        0/4     Terminating    0          5m

student@lfs458-node-1a0a:~$ kubectl get pods
No resources found.

student@lfs458-node-1a0a:~$ kubectl create -f vip.yaml
pod "vip" created

```

10. Determine where the new containers have been deployed.

```

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
19

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
12

```

11. Create another file for other users. Change the names from vip to others, and uncomment the nodeSelector lines.

```

student@lfs458-node-1a0a:~$ cp vip.yaml other.yaml

student@lfs458-node-1a0a:~$ sed -i s/vip/other/g other.yaml

student@lfs458-node-1a0a:~$ vim other.yaml
.
nodeSelector:
  status: other

```

12. Create the other containers. Determine where they deploy.

```
student@lfs458-node-1a0a:~$ kubectl create -f other.yaml
pod "other" created

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
19

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
17
```

13. Shut down both pods and verify they are terminating.

```
student@lfs458-node-1a0a:~$ kubectl delete pods vip other
pod "vip" deleted
pod "other" deleted

student@lfs458-node-1a0a:~$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
other     4/4     Terminating    0          5m
vip       4/4     Terminating    0          10m
```

Exercise 11.2: Using Taints to Control Pod Deployment

Use taints to manage where Pods are deployed or allowed to run. In addition to assigning a Pod to a group of nodes, you may also want to limit usage on a node or fully evacuate Pods. Using taints is one way to achieve this. You may remember that the master node begins with a `NoSchedule` taint. We will work with three taints to limit or remove running pods.

1. Verify that the master and secondary node have the minimal number of containers running.
2. Create a deployment which will deploy eight **nginx** containers. Begin by creating a YAML file.

```
student@lfs458-node-1a0a:~$ vim taint.yaml
```

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: taint-deployment
spec:
  replicas: 8
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
```

3. Apply the file to create the deployment.

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment "taint-deployment" created
```

4. Determine where the containers are running. In the following example three have been deployed on the master node and five on the secondary node. Remember there will be other housekeeping containers created as well.

```
student@lfs458-node-1a0a:~$ sudo docker ps |grep nginx
00c1be5df1e7      nginx@sha256:e3456c851a152494c3e..
<output_omitted>
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
25
student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
17
```

5. Delete the deployment. Verify the containers are gone.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment "taint-deployment" deleted
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
19
```

6. Now we will use a taint to affect the deployment of new containers. There are three taints, `NoSchedule`, `PreferNoSchedule` and `NoExecute`. The taints having to do with schedules will be used to determine newly deployed containers, but will not affect running containers. The use of `NoExecute` will cause running containers to move.

Taint the secondary node, verify it has the taint then create the deployment again. We will use the key of `bubba` to illustrate the key name is just some string an admin can use to track Pods.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-node-2b2b bubba=value:PreferNoSchedule
node "lfs458-node-2b2b" tainted
```

```
student@lfs458-node-1a0a:~$ kubectl describe node |grep Taint
Taints:          bubba=value:PreferNoSchedule
Taints:          <none>
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment "taint-deployment" created
```

7. Locate where the containers are running. We can see that more containers are on the master, but there still were some created on the secondary. Delete the deployment when you have gathered the numbers.

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
29
```

```
student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
13
```

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment "taint-deployment" deleted
```

8. Remove the taint, verify it has been removed. Note that the key is used with a minus sign appended to the end.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-node-2b2b bubba-
node "lfs458-node-2b2b" untainted
```

```
student@lfs458-node-1a0a:~$ kubectl describe node |grep Taint
Taints:          <none>
Taints:          <none>
```

9. This time use the `NoSchedule` taint, then create the deployment again. The secondary node should not have any new containers.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-node-2b2b bubba=value:NoSchedule
node "lfs458-node-2b2b" tainted
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
```

```
deployment "taint-deployment" created

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
35

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
7
```

10. Remove the taint and delete the deployment. When you have determined that all the containers are running create the deployment again. Without any taint the containers should be spread across both nodes.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment "taint-deployment" deleted

student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-node-2b2b bubba-
node "lfs458-node-2b2b" untainted

student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment "taint-deployment" created

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
25

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
17
```

11. Now use the NoExecute to taint the secondary node. Wait a minute then determine if the containers have moved. The DNS containers can take a while to shutdown.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-node-2b2b bubba=value:NoExecute
node "lfs458-node-2b2b" tainted
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
35

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
1
```

12. Remove the taint. Wait a minute. Note that all of the containers did not return to their previous placement.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-node-2b2b bubba-
node "lfs458-node-2b2b" untainted

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
35

student@lfs458-node-2b2b:~$ sudo docker ps |wc -l
7
```

13. In addition to the ability to taint a node you can also set the status to drain. First view the status, then destroy the existing deployment. Note that the status reports Ready, even though it will not allow containers to be executed. Also note that the output mentioned that DaemonSet-managed pods are not affected by default.

In v1.9.1 there appears to be an issue where the existing containers are not moved, but no new containers are created. You may receive an error error: unable to drain node "<your node>", aborting command...

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    AGE      VERSION
lfs458-node-1a0a    Ready     2d       v1.9.1
lfs458-node-2b2b    Ready     <none>    2d       v1.9.1

student@lfs458-node-1a0a:~$ kubectl drain lfs458-node-2b2b
node "lfs458-node-2b2b" cordoned
error: DaemonSet-managed pods (use --ignore-daemonsets to ignore): kube-flannel-ds-fx3tx, kube-proxy-q2q4k
```

14. Verify the state change of the node. It should indicate no new Pods will be scheduled.

```
student@lfs458-node-1a0a:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
lfs458-node-1a0a	Ready	master	2d	v1.9.1
lfs458-node-2b2b	Ready,SchedulingDisabled	<none>	2d	v1.9.1

15. Delete the deployment to destroy the current Pods.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment "taint-deployment" deleted
```

16. Create the deployment again and determine where the containers have been deployed.

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment "taint-deployment" created

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
35
```

17. Return the status to Ready, then destroy and create the deployment again. The containers should be spread across the nodes. Begin by removing the cordon on the node.

```
student@lfs458-node-1a0a:~$ kubectl uncordon lfs458-node-2b2b
node "lfs458-node-2b2b" uncordoned

student@lfs458-node-1a0a:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
lfs458-node-1a0a	Ready	master	2d	v1.9.1
lfs458-node-2b2b	Ready	<none>	2d	v1.9.1

18. Delete and re-create the deployment.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment "taint-deployment" deleted

student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment "taint-deployment" created
```

19. View the **docker ps** output again. Both nodes should have almost the same number of containers deployed. The master will have a few more, due to its role.

20. Remove the deployment a final time to free up resources.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
```


Chapter 12

Logging and Troubleshooting



12.1 Labs

Exercise 12.1: Review Log File Locations

Overview

In addition to various logs files and command output, you can use **journalctl** to view logs from the node perspective. We will view common locations of log files, then a command to view container logs. There are other logging options, such as the use of a **sidecar** container dedicated to loading the logs of another container in a pod.

Whole cluster logging is not yet available with Kubernetes. Outside software is typically used, such as **Fluentd**, part of <https://fluentd.org/>, which is another member project of **CNCF**, like Kubernetes.

exe

Review Log File Locations

Take a quick look at the following log files and web sites. As server processes move from node level to running in containers the logging also moves.

1. If using a **systemd** based Kubernetes cluster view the node level logs for **kubelet**, the local Kubernetes agent. Each node will have different contents as this is node specific.

```
student@lfs458-node-1a0a:~$ journalctl -u kubelet |less
<output_omitted>
```

2. Major Kubernetes processes now run in containers. You can view them from the container or the pod perspective. Use the **find** command to locate the **kube-apiserver** log. Your output will be different, but will be very long. Once you locate the files use the **diff** utility to compare them. There should be no difference.

```
student@lfs458-node-1a0a:~$ sudo find / -name "*apiserver*log"
/var/log/pods/56c55117e68ed986eaddeb0f78ca405e/kube-apiserver_0.log
/var/log/containers/kube-apiserver-lfs458-node-9q6r_kube-system_kube...
```

```
student@lfs458-node-1a0a:~$ sudo diff \
/var/log/pods/56c55117e68ed986eaddeb0f78ca405e/kube-apiserver_0.log
/var/log/containers/kube-apiserver-lfs458-node-9q6r_kube-system_kube...
```

3. Take a look at the log file.

```
student@lfs458-node-1a0a:~$ sudo less \
/var/log/pods/56c55117e68ed986eaddeb0f78ca405e/kube-apiserver_0.log
```

4. Search for and review other log files for kube-dns, kube-flannel, and kube-proxy.
5. If not on a Kubernetes cluster using **systemd** you can view the text files on the master node.

- (a) `/var/log/kube-apiserver.log`
Responsible for serving the API
- (b) `/var/log/kube-scheduler.log`
Responsible for making scheduling decisions
- (c) `/var/log/kube-controller-manager.log`
Controller that manages replication controllers

6. `/var/log/containers`
Various container logs

7. `/var/log/pods/`
More log files for current Pods.

8. Worker Nodes Files (on non-**systemd** systems)

- (a) `/var/log/kubelet.log`
Responsible for running containers on the node
- (b) `/var/log/kube-proxy.log`
Responsible for service load balancing

9. More reading: <https://kubernetes.io/docs/tasks/debug-application-cluster/\debug-service/> and <https://kubernetes.io/docs/tasks/debug-application-cluster/\determine-reason-pod-failure/>

Exercise 12.2: Viewing Logs Output

Container standard out can be seen via the **kubectl logs** command. If there is no standard out, you would not see any output. In addition, the logs would be destroyed if the container is destroyed.

1. View the current Pods in the cluster. Be sure to view Pods in all namespaces.

```
student@lfs458-node-1a0a:~$ kubectl get po --all-namespaces
NAMESPACE   NAME                                READY   STATUS    RESTARTS   AGE
default     ds-one-qc72k                       1/1     Running   0           3h
default     ds-one-z31r4                       1/1     Running   0           3h
...
kube-system etcd-lfs458-node-1a0a        1/1     Running   2           9h
kube-system kube-apiserver-lfs458-node-1a0a  1/1     Running   2           9h
kube-system kube-controller-manager-lfs458-node-1a0a  1/1     Running   2           9h
kube-system kube-dns-2425271678-w80vx          3/3     Running   6           9h
kube-system kube-scheduler-lfs458-node-1a0a    1/1     Running   2           9h
...
```

2. View the logs associated with the container scheduler. Your particular Pod name will be different. You will get a lot of output. Consider narrowing down the output to only the **nginx** server.

```
student@lfs458-node-1a0a:~$ kubectl logs kube-scheduler-lfs458-node-1a0a\
-n kube-system | grep nginx
```

```
....
```

```
I1024 19:47:54.955077      1 event.go:218] Event(v1.ObjectReference
{Kind:"Pod", Namespace:"default", Name:"nginx-2060361664-4g2g2",
UID:"3a670a6f-b8f4-11e7-8350-0a053b35c9b8", APIVersion:"v1",
ResourceVersion:"809652", FieldPath:""}): type: 'Normal' reason:
'Scheduled' Successfully assigned nginx-2060361664-4g2g2 to
lfs458-node-1a0a
```

```
I1024 19:49:33.096360      1 event.go:218] Event(v1.ObjectReference
{Kind:"Pod", Namespace:"default", Name:"nginx-2060361664-zbv36",
UID:"74e66be7-b8f4-11e7-8350-0a053b35c9b8", APIVersion:"v1",
ResourceVersion:"809834", FieldPath:""}): type: 'Normal'
reason: 'Scheduled' Successfully assigned nginx-2060361664-zbv36
to lfs458-node-1a0a
```

3. View the logs of other Pods in your cluster.

Chapter 13

Custom Resource Definition



13.1 Labs

Exercise 13.1: Create a Custom Resource Definition

Overview

ThirdPartyResource is no longer included with the API in v1.8 and its use will return a validation error. If you have upgraded from a version prior to Kubernetes v1.7, you will need to convert them to CustomResourceDefinitions (CRD). A new resource often requires a controller to manage the resource. Creation of the controller is beyond the scope of this course, basically it is a watch-loop comparing a spec file to the current state and making changes until the states match. A good discussion of creating a controller can be found here: <https://coreos.com/blog/introducing-operators.html>.

Create a Custom Resource Definition

We will make a simple CRD, but without any particular action. It will be enough to find the object ingested into the API and responding to commands.

1. We will create a new YAML file.

```
student@lfs458-node-1a0a:~$ vim crd.yaml
```

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: crontabs.training.lfs458.com
    # This name must match names below.
    # <plural>.<group> syntax
spec:
  scope: Cluster      #Could also be Namespaced
  group: training.lfs458.com
```

```

version: v1
names:
  kind: CronTab      #Typically CamelCased for resource manifest
  plural: crontabs   #Shown in URL
  singular: crontab  #Short name for CLI alias
  shortNames:
    - ct             #CLI short name

```

2. Add the new resource to the cluster.

```

student@lfs458-node-1a0a:~$ kubectl create -f crd.yaml
customresourcedefinition "crontabs.training.lfs458.com" created

```

3. View and describe the resource.

```

student@lfs458-node-1a0a:~$ kubectl get crd
NAME                                AGE
crontabs.training.lfs458.com        3m

student@lfs458-node-1a0a:~$ kubectl describe crd
Name:          crontabs.training.lfs458.com
Namespace:
Labels:        <none>
Annotations:   <none>
API Version:   apiextensions.k8s.io/v1beta1
Kind:          CustomResourceDefinition
<output_omitted>

```

4. Now that we have a new API resource we can create a new object of that type. In this case it will be a crontab-like image, which does not actually exist, but is being used for demonstration.

```

student@lfs458-node-1a0a:~$ vim new-crontab.yaml

apiVersion: "training.lfs458.com/v1"
  # This is from the group and version of new CRD
kind: CronTab
  # The kind from the new CRD
metadata:
  name: new-cron-object
spec:
  cronSpec: "*/5 * * * *"
  image: some-cron-image
  #Does not exist

```

5. Create the new object and view the resource using short and long name.

```

student@lfs458-node-1a0a:~$ kubectl create -f new-crontab.yaml
crontab "new-cron-object" created

student@lfs458-node-1a0a:~$ kubectl get CronTab
NAME                                AGE
new-cron-object                    22s

student@lfs458-node-1a0a:~$ kubectl get ct
NAME                                AGE
new-cron-object                    29s

student@lfs458-node-1a0a:~$ kubectl describe ct
Name:          new-cron-object
Namespace:
Labels:        <none>

```

```
<output_omitted>
```

```
Spec:
```

```
  Cron Spec: */5 * * * *
```

```
  Image:      some-cron-image
```

```
Events:      <none>
```

6. To clean up the resources we will delete the CRD. This should delete all of the endpoints and objects using it as well.

```
student@lfs458-node-1a0a:~$ kubectl delete -f crd.yaml
customresourcedefinition "crontabs.training.lfs458.com" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get ct
Error from server (NotFound): Unable to list "crontabs": the server
could not find the requested resource
(get crontabs.training.lfs458.com)
```


Chapter 14

Kubernetes Federation



14.1 Labs

There is no lab to complete for this chapter.

Chapter 15

Helm



15.1 Labs

Exercise 15.1: Working with Helm and Charts

Overview

helm allows for easy deployment of complex configurations. This could be handy for a vendor to deploy a multi-part application in a single step. Through the use of a Chart, or template file, the required components and their relationships are declared. Local agents like **Tiller** use the API to create objects on your behalf. Effectively its orchestration for orchestration.

There are a few ways to install **Helm**. The newest version may require building from source code. We will download a recent, stable version. Once installed we will deploy a Chart, which will configure **Hadoop** on our cluster.

Install Helm

1. On the master node use **wget** to download the compressed tar file. The short URL below is for: <https://storage.googleapis.com/kubernetes-helm/helm-v2.7.0-linux-amd64.tar.gz>

```
student@lfs458-node-1a0a:~$ wget goo.gl/nbEcHn
```

2. Uncompress and expand the file.

```
student@lfs458-node-1a0a:~$ tar -xvf nbEcHn
linux-amd64/
linux-amd64/README.md
linux-amd64/helm
linux-amd64/LICENSE
```

3. Copy the **helm** binary to the `/usr/local/bin/` directory, so it is usable via the shell search path.

```
student@lfs458-node-1a0a:~$ sudo cp linux-amd64/helm /usr/local/bin/
```

- Due to new **RBAC** configuration **helm** is unable to run in the default namespace, in this version of Kubernetes. During initialization you could choose to create and declare a new namespace. Other RBAC issues may be encountered even then. In this lab we will create a service account for **tiller**, and give it admin abilities on the cluster. More on **RBAC** in another chapter.

Begin by creating the serviceaccount object.

```
student@lfs458-node-1a0a:~$ kubectl create serviceaccount \
--namespace kube-system tiller
serviceaccount "tiller" created
```

- Bind the serviceaccount to the admin role called cluster-admin inside the kube-system namespace.

```
student@lfs458-node-1a0a:~$ kubectl create clusterrolebinding \
tiller-cluster-rule \
--clusterrole=cluster-admin \
--serviceaccount=kube-system:tiller
clusterrolebinding "tiller-cluster-rule" created
```

- Update the tiller-deploy deployment to have the service account.

```
student@lfs458-node-1a0a:~$ kubectl patch deployment \
tiller-deploy -p \
'{"spec":{"template":{"spec":{"serviceAccount":"tiller"}}}}' \
-n kube-system
```

- We can now initialize **helm**. This process will also configure **tiller** the client process. There are several possible options to pass such as nodeAffinity, a particular version of software, alternate storage backend, and even a dry-run option to generate JSON or YAML output. The output could be edited and ingested into **kubectl**. We will use default values in this case.

```
student@lfs458-node-1a0a:~$ helm init
```

```
<output_omitted>
```

- Verify the **tiller** pod is running. Examine the logs of the pod. Note that each line of log begins with an tag of the component generating the messages, such as [main], [storage], and [storage].

```
student@lfs458-node-1a0a:~$ kubectl get pods --all-namespaces
<output_omitted>
kube-system    tiller-deploy-84b97f465c-76lvs    1/1    Running    0    30m
```

```
student@lfs458-node-1a0a:~$ kubectl logs tiller-deploy-84b97f465c-76lvs\
-n kube-system
<output_omitted>
```

- View the available sub-commands for **helm**. As with other Kubernetes tools, expect ongoing change.

```
student@lfs458-node-1a0a:~$ helm help
<output_omitted>
```

- View the current configuration files, archives and plugins for helm. Return to this directory after you have worked with a Chart later in the lab.

```
student@lfs458-node-1a0a:~$ helm home
/home/student/.helm
```

```
student@lfs458-node-1a0a:~$ ls -R /home/student/.helm/
/home/student/.helm/:
cache  plugins  repository  starters
```

```
/home/student/.helm/cache:
archive
<output_omitted>
```

11. Verify **helm** and **tiller** are responding, also check the current version installed.

```
student@lfs458-node-1a0a:~$ helm version
Client: &version.Version{SemVer:"v2.7.0", GitCommit:"08c1144f5...
Server: &version.Version{SemVer:"v2.7.0", GitCommit:"08c1144f5...
```

12. Ensure both are upgraded to the most recent stable version.

```
student@lfs458-node-1a0a:~$ helm init --upgrade
```

13. A Chart is a collection of containers to deploy an application. There is a collection available on <https://github.com/kubernetes/charts/tree/master/stable>, provided by vendors, or you can make your own. Take a moment and view the current stable Charts. Then search for available stable databases.

```
student@lfs458-node-1a0a:~$ helm search database
NAME                VERSION  DESCRIPTION
stable/cockroachdb  0.5.4    CockroachDB is a scalable, ...
stable/dokuwiki     0.2.1    DokuWiki is a standards-compliant...
stable/mariadb      2.1.3    Fast, reliable, scalable, and eas...
stable/mediawiki    0.6.1    Extremely powerful, scalable soft...
stable/mongodb      0.4.22   NoSQL document-oriented database ...
stable/mongodb-replicaset 2.1.4    NoSQL document-oriented datab...
stable/mysql        0.3.4    Fast, reliable, scalable, and eas...
<output_omitted>
```

14. We will install the **mariadb**. Take a look at install details <https://github.com/kubernetes/charts/tree/master/stable/mariadb#custom-mycnf-configuration> The **-debug** option will create a lot of output. Note the interesting name for the deployment, like **illmannered-salamander**. The output will typically suggest ways to access the software. As well we will indicate that we do not want persistent storage, which would require use to create an available PV.

```
student@lfs458-node-1a0a:~$ helm --debug install stable/mariadb \
--set persistence.enabled=false
[debug] Created tunnel using local port: '38396'

[debug] SERVER: "localhost:38396"

[debug] Original chart version: ""
[debug] Fetched stable/mariadb to /home/student/.helm/cache/archive/mar...

[debug] CHART PATH: /home/student/.helm/cache/archive/mariadb-2.0.1.tgz

NAME:    illmannered-salamander
<output_omitted>
```

15. Using some of the information at the end of the previous command output we will deploy another container and access the database. We begin by getting the root password for **illmannered-salamander**. Be aware the output lacks a carriage return, so the next prompt will appear on the same line. We will need the password to access the running MariaDB database.

```
student@lfs458-node-1a0a:~$ kubectl get secret --namespace default \
illmannered-salamander-mariadb \
-o jsonpath="{.data.mariadb-root-password}" \
| base64 --decode
IFBldzAQfx
```

16. Now we will install another container to act as a client for the database. We will use **apt-get** to install client software.

```
student@lfs458-node-1a0a:~$ kubectl run -i --tty ubuntu --image=ubuntu:16.04 --restart=Never -- bash -il
If you don't see a command prompt, try pressing enter.
root@ubuntu:/#

root@ubuntu:/# apt-get update ; apt-get install -y mariadb-client
<output_omitted>
```

17. Use the client software to access the database. The following command uses the server name and the root password we found in a previous step. Both of yours will be different.

```
root@ubuntu:/# mysql -h illmannered-salamander-mariadb -p
Enter password: IFBldzAQfx
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MariaDB connection id is 153
Server version: 10.1.28-MariaDB Source distribution

Copyright (c) 2000, 2017, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [(none)]> SHOW DATABASES;
+-----+
| Database                |
+-----+
| information_schema      |
| mysql                   |
| performance_schema      |
+-----+
3 rows in set (0.00 sec)

MariaDB [(none)]>
MariaDB [(none)]> quit
root@ubuntu:/# exit
```

18. View the Chart history on the system. The use of the **-a** option will show all Charts including deleted and failed attempts. The output below shows the current running Chart as well as a previously deleted **hadoop** Chart.

```
student@lfs458-node-1a0a:~$ helm list -a
NAME            REVISION    UPDATED      STATUS  CHART          NAMESPACE
goodly-beetle   1           Wed Nov  8 23:01:24 2017 DELETED hadoop-1.0.1 default
illmannered-salamander 1           Thu Nov  9 05:00:12 2017 DEPLOYED mariadb-...
```

19. Delete the **mariadb** Chart. No output should happen from the list.

```
student@lfs458-node-1a0a:~$ helm delete illmannered-salamander
release "illmannered-salamander" deleted

student@lfs458-node-1a0a:~$ helm list
```

20. Add another repository and view the Charts available.

```
student@lfs458-node-1a0a:~$ helm repo add stable \
    http://storage.googleapis.com/kubernetes-charts

student@lfs458-node-1a0a:~$ helm search
NAME                                VERSION    DESCRIPTION
stable/acs-engine-autoscaler        2.1.0      Scales worker nodes within...
stable/artifactory                   6.2.0      Universal Repository Manag...
<output_omitted>
```

Chapter 16

Security



16.1 Labs

Exercise 16.1: Working with TLS

Overview

We have learned that the flow of access to a cluster begins with TLS connectivity, then authentication followed by authorization, finally an admission control plug-in allows advanced features prior to the request being fulfilled. The use of `Initializers` allows the flexibility of a shell-script to dynamically modify the request. As security is an important, ongoing concern there may be multiple configurations used depending on the needs of the cluster.

Every process making API requests to the cluster must authenticate or be treated as an anonymous user.

Working with TLS

While one can have multiple cluster root Certificate Authorities (CA) by default each cluster uses their own, intended for intra-cluster communication. The CA certificate bundle is distributed to each node and as a secret to default service accounts. The **kubelet** is a local agent which ensures local containers are running and healthy.

1. View the **kubelet** on both the master and secondary nodes. The **kube-apiserver** also shows security information such as certificates and authorization mode.

```
student@lfs458-node-1a0a:~$ ps -ef |grep kubelet
<output_omitted>
--cluster-domain=cluster.local --authorization-mode=Webhook
--client-ca-file=/etc/kubernetes/pki/ca.crt --cadvisor-port=0
--rotate-certificates=true --cert-dir=/var/lib/kubelet/pki
<output_omitted>
--service-account-key-file=/etc/kubernetes/pki/sa.pub --secure-port=6443
--requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt
--proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt
```

```

<output_omitted>
--authorization-mode=Node,RBAC
--etcd-servers=http://127.0.0.1:2379

student@lfs458-node-2b2b:~$ ps -ef |grep kubelet
<output_omitted>
--authorization-mode=Webhook --client-ca-file=/etc/kubernetes/pki/ca.crt
--cadvisor-port=0 --rotate-certificates=true
--cert-dir=/var/lib/kubelet/pki

```

2. View the configuration file where these settings are made.

```

student@lfs458-node-1a0a:~$ ls /etc/kubernetes/manifests/
etcd.yaml                kube-controller-manager.yaml
kube-apiserver.yaml      kube-scheduler.yaml

student@lfs458-node-1a0a:~$ sudo less /etc/kubernetes/manifests/kube-controller-manager.yaml
<output_omitted>

```

3. Note the certificate used is the same as that from `~/.kube/config`.

```

student@lfs458-node-1a0a:~$ kubectl get csr -o yaml
<output_omitted>
  status:
    certificate: LS0tLS1CRUdJTiBDRVJUSUZJ.....
<output_omitted>

student@lfs458-node-1a0a:~$ grep cert ~/.kube/config
<output_omitted>

```

4. The **kubectl config** command can also be used to update parameters, which could avoid a typo removing access to the cluster. View the current configuration settings.

```

student@lfs458-node-1a0a:~$ kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
<output_omitted>

```

5. View the options, such as setting a password for the admin instead of a key. Read through the examples and options.

```

student@lfs458-node-1a0a:~$ kubectl config set-credentials -h
Sets a user entry in kubeconfig
<output_omitted>

```

6. Make a copy of your access configuration file. Later steps will update this file and we can view the differences.

```

student@lfs458-node-1a0a:~$ cp ~/.kube/config ~/cluster-api-config

```

7. Explore working with cluster and security configurations both using **kubectl** and **kubeadm**. Among other values, find the name of your cluster. You will need to become root to work with **kubeadm**.

```

student@lfs458-node-1a0a:~$ kubectl config <Tab><Tab>
current-context  get-contexts      set-context      view
delete-cluster   rename-context    set-credentials
delete-context   set               unset
get-clusters     set-cluster       use-context

student@lfs458-node-1a0a:~$ sudo -i

root@lfs458-node-1a0a:~# kubeadm token -h
<output_omitted>

```



```
root@lfs458-node-1a0a:~# kubeadm config -h
<output_omitted>
```

Exercise 16.2: Authentication and Authorization

Kubernetes clusters have two types of users: service accounts and normal users, but normal users are assumed to be managed by an outside service. There are no objects to represent them and they cannot be added via an API call, but service accounts can be added.

We will use **RBAC** to configure access to actions within a namespace for a new contractor, Developer Dan who will be working on a new project.

1. Create two namespaces, one for production and the other for development.

```
student@lfs458-node-1a0a:~$ kubectl create ns development
namespace "development" created

student@lfs458-node-1a0a:~$ kubectl create ns production
namespace "production" created
```

2. View the current clusters and context available. The context allows you to configure the cluster to use, namespace and user for **kubectl** commands in an easy and consistent manner.

```
student@lfs458-node-1a0a:~$ kubectl config get-contexts
CURRENT  NAME                      CLUSTER          AUTHINFO          NAMESPACE
*        kubernet-admin@kubernetes  kubernet          kubernet          kubernet-admin
```

3. Create a new user DevDan and assign a password of lfs458.

```
student@lfs458-node-1a0a:~$ sudo useradd -s /bin/bash DevDan
student@lfs458-node-1a0a:~$ sudo passwd DevDan
Enter new UNIX password: lfs458
Retype new UNIX password: lfs458
passwd: password updated successfully
```

4. Generate a private key then Certificate Signing Request (CSR) for DevDan.

```
student@lfs458-node-1a0a:~$ openssl genrsa -out DevDan.key 2048
Generating RSA private key, 2048 bit long modulus
.....+++
.....+++
e is 65537 (0x10001)

student@lfs458-node-1a0a:~$ openssl req -new -key DevDan.key \
-out DevDan.csr -subj "/CN=DevDan/O=development"
```

5. Using the newly created request generate a self-signed certificate using the x509 protocol. Use the CA keys for the Kubernetes cluster and set a 45 day expiration. You'll need to use **sudo** to access to the inbound files.

```
student@lfs458-node-1a0a:~$ sudo openssl x509 -req -in DevDan.csr \
-CA /etc/kubernetes/pki/ca.crt \
-CAkey /etc/kubernetes/pki/ca.key \
-CACreateserial \
-out DevDan.crt -days 45
Signature ok
subject=/CN=DevDan/O=development
Getting CA Private Key
```

6. Update the access config file to reference the new key and certificate. Normally we would move them to a safe directory instead of a non-root user's home.

```
student@lfs458-node-1a0a:~$ kubectl config set-credentials DevDan \
--client-certificate=/home/student/DevDan.crt \
--client-key=/home/student/DevDan.key
User "DevDan" set.
```

7. View the update to your credentials file. Use **diff** to compare against the copy we made earlier.

```
student@lfs458-node-1a0a:~$ diff cluster-api-config .kube/config
9a10,14
> namespace: development
> user: DevDan
> name: DevDan-context
> - context:
>   cluster: kubernetes
15a21,25
> - name: DevDan
>   user:
>   as-user-extra: {}
>   client-certificate: /home/tim/DevDan.crt
>   client-key: /home/tim/DevDan.key
```

8. We will now create a context. For this we will need the name of the cluster, namespace and CN of the user we set or saw in previous steps.

```
student@lfs458-node-1a0a:~$ kubectl config set-context DevDan-context \
--cluster=kubernetes \
--namespace=development \
--user=DevDan
Context "DevDan-context" created.
```

9. Attempt to view the Pods inside the DevDan-context. Be aware you will get an error.

```
student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context get pods
Error from server (Forbidden): pods is forbidden: User "DevDan"
cannot list pods in the namespace "development"
```

10. Verify the context has been properly set.

```
student@lfs458-node-1a0a:~$ kubectl config get-contexts
CURRENT  NAME             CLUSTER      AUTHINFO      NAMESPACE
         DevDan-context   kubernetes   DevDan        development
*        kubernetes-admin@kubernetes kubernetes   kubernetes-admin
```

11. Again check the recent changes to the cluster access config file.

```
student@lfs458-node-1a0a:~$ diff cluster-api-config .kube/config
9a10,14
> namespace: development
> user: DevDan
> name: DevDan-context
> - context:
>   cluster: kubernetes
15a21,25
> - name: DevDan
>   user:
<output_omitted>
```

12. We will now create a YAML file to associate RBAC rights to a particular namespace and Role.

```
student@lfs458-node-1a0a:~$ vim role-dev.yaml
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: development
```

```

    name: developer
rules:
- apiGroups: ["", "extensions", "apps"]
  resources: ["deployments", "replicasets", "pods"]
  verbs: ["list", "get", "watch", "create", "update", "patch", "delete"]
# You can use ["*"] for all verbs

```

13. Create the object. Check white space and for typos if you encounter errors.

```

student@lfs458-node-1a0a:~$ kubectl create -f role-dev.yaml
role "developer" created

```

14. Now we create a RoleBinding to associate the Role we just created with a user. Create the object when the file has been created.

```

student@lfs458-node-1a0a:~$ vim rolebind.yaml

```

```

kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: developer-role-binding
  namespace: development
subjects:
- kind: User
  name: DevDan
  apiGroup: ""
roleRef:
  kind: Role
  name: developer
  apiGroup: ""

```

```

student@lfs458-node-1a0a:~$ kubectl apply -f rolebind.yaml
rolebinding "developer-role-binding" created

```

15. Test the context again. This time it should work. There are no Pods running so you should get a response of No resources found.

```

student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context get pods
No resources found.

```

16. Create a new pod, verify it exists, then delete it.

```

student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context run nginx --image=nginx
deployment "nginx" created

```

```

student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-7c87f569d-7gb9k              1/1     Running   0           5s

```

```

student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context delete deploy nginx
deployment "nginx" deleted

```

17. We will now create a different context for production systems. The Role will only have the ability to view, but not create or delete resources. Begin by copying and editing the Role and RoleBindings YAML files.

```

student@lfs458-node-1a0a:~$ cp role-dev.yaml role-prod.yaml

```

```

student@lfs458-node-1a0a:~$ vim role-prod.yaml

```

```

kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:

```

```

    namespace: production      #<<- This line
    name: dev-prod             #<<- and this line
rules:
- apiGroups: ["", "extensions", "apps"]
  resources: ["deployments", "replicasets", "pods"]
  verbs: ["get", "list", "watch"] #<<- and this one

student@lfs458-node-1a0a:~$ cp rolebind.yaml rolebindprod.yaml

student@lfs458-node-1a0a:~$ vim rolebindprod.yaml

kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: production-role-binding
  namespace: production
subjects:
- kind: User
  name: DevDan
  apiGroup: ""
roleRef:
  kind: Role
  name: dev-prod
  apiGroup: ""

```

18. Create both new objects.

```

student@lfs458-node-1a0a:~$ kubectl apply -f role-prod.yaml
role "dev-prod" created

student@lfs458-node-1a0a:~$ kubectl apply -f rolebindprod.yaml
rolebinding "production-role-binding" created

```

19. Create the new context for production use.

```

student@lfs458-node-1a0a:~$ kubectl config set-context ProdDan-context \
    --cluster=kubernetes \
    --namespace=production \
    --user=DevDan
Context "ProdDan-context" created.

```

20. Verify that user DevDan can view pods using the new context.

```

student@lfs458-node-1a0a:~$ kubectl --context=ProdDan-context get pods
No resources found.

```

21. Try to create a Pod in production. The developer should be Forbidden.

```

student@lfs458-node-1a0a:~$ kubectl --context=ProdDan-context run \
    nginx --image=nginx
Error from server (Forbidden): deployments.extensions is forbidden: User "DevDan" cannot create deployments.extensions

```

22. View the details of a role.

```

student@lfs458-node-1a0a:~$ kubectl describe role dev-prod -n production
Name:          dev-prod
Labels:        <none>
Annotations:   kubectl.kubernetes.io/last-applied-configuration=
{"apiVersion":"rbac.authorization.k8s.io/v1beta1","kind":"Role",
"metadata":{"annotations":{"name":"dev-prod","namespace":
"production"},"rules":[{"api...
PolicyRule:
  Resources      Non-Resource URLs  Resource Names  Verbs

```

```

-----
deployments      []          []          [get list watch]
deployments.apps []          []          [get list watch]
<output_omitted>

```

23. Experiment with other subcommands in both contexts. They should match those listed in the respective roles.

Exercise 16.3: Admission Controllers

The last stop before a request is sent to the API server is an admission control plug-in. They interact with features such as setting parameters like a default storage class, checking resource quotas, or security settings. A newer feature (v1.7.x) is dynamic controllers which allow new controllers to be ingested or configured at runtime.

The order of listing matters. When an admission controller matches the plug-ins action is performed, acceptance or mutation of the request, and the request is sent along to the API server. No further plug-ins are checked.

1. View the current admission controller list and order. Note that Initializers, a dynamic controller, is listed first. This allows modification of the object. The use of webhooks could also be used if notification is the only desired goal.

```

student@lfs458-node-1a0a:~$ sudo grep admission \
    /etc/kubernetes/manifests/kube-apiserver.yaml
- --admission-control=Initializers,NamespaceLifecycle,
  LimitRanger,ServiceAccount,PersistentVolumeLabel,
  DefaultStorageClass,DefaultTolerationSeconds,NodeRestriction,
  ResourceQuota

```